

# CSCI 6760/4760 Computer Networks: Topology and applications

Chapter 2

Manijeh Keshtgari

Computer Science  
UGA  
[M.keshtgari@uga.edu](mailto:M.keshtgari@uga.edu)

# Chapter 2: application layer

## our goals:

- ❖ conceptual, implementation aspects of network application protocols
  - transport-layer service models
  - client-server paradigm
  - peer-to-peer paradigm
- ❖ learn about protocols by examining popular application-level protocols
  - HTTP
  - FTP
  - SMTP / POP3 / IMAP
  - DNS
- ❖ creating network applications
  - socket API

# Some network apps

- ❖ e-mail
- ❖ web
- ❖ text messaging
- ❖ remote login
- ❖ P2P file sharing
- ❖ multi-user network games
- ❖ streaming stored video  
(YouTube, Hulu, Netflix)
- ❖ voice over IP (e.g., Skype)
- ❖ real-time video conferencing
- ❖ social networking
- ❖ search
- ❖ ...
- ❖ ...

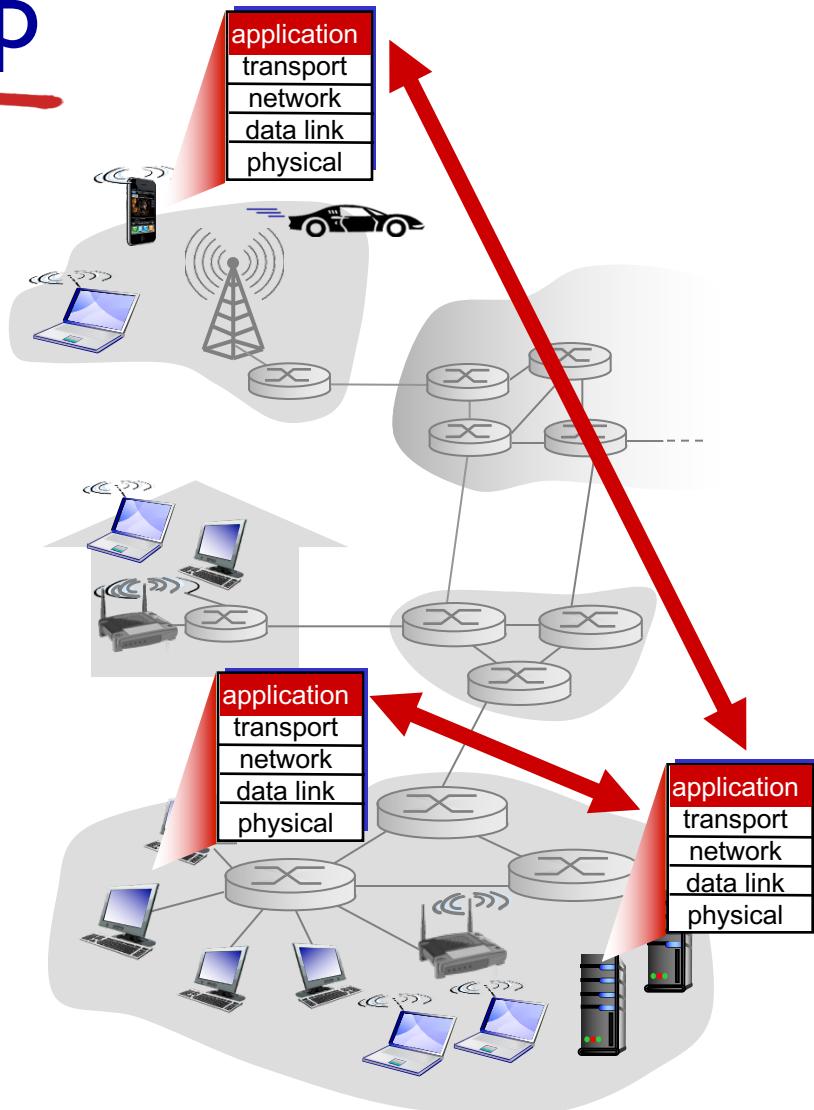
# Creating a network app

write programs that:

- ❖ run on (different) end systems
- ❖ communicate over network
- ❖ e.g., web server software communicates with browser software

no need to write software for network-core devices

- ❖ network-core devices do not run user applications
- ❖ applications on end systems allows for rapid app development, propagation

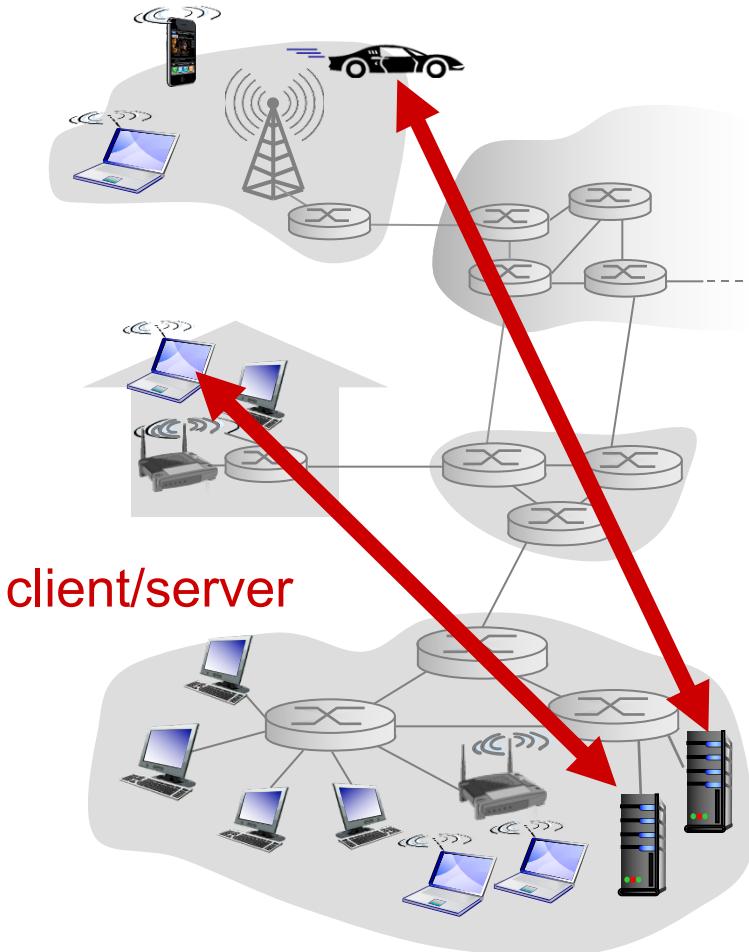


# Application architectures

possible structure of applications:

- ❖ client-server
- ❖ peer-to-peer (P2P)

# Client-server architecture



## server:

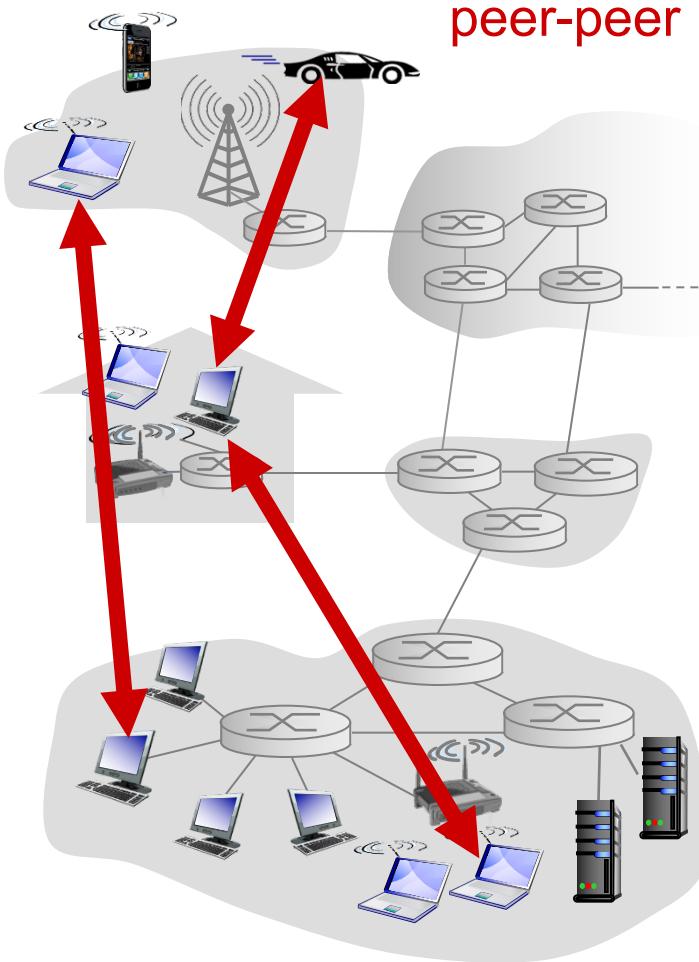
- ❖ always-on host
- ❖ permanent IP address
- ❖ data centers for scaling

## clients:

- ❖ communicate with server
- ❖ may have dynamic IP addresses
- ❖ do not communicate directly with each other

# P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers request service from other peers, provide service in return to other peers
  - *self scalability* – new peers bring new service capacity, as well as new service demands
- ❖ Example: BitTorrent
- ❖ peers are intermittently connected and change IP addresses
  - complex management



# Processes communicating

*process*: program running within a host

- ❖ within same host, two processes communicate using **inter-process communication** (defined by OS)
- ❖ processes in different hosts communicate by exchanging **messages**

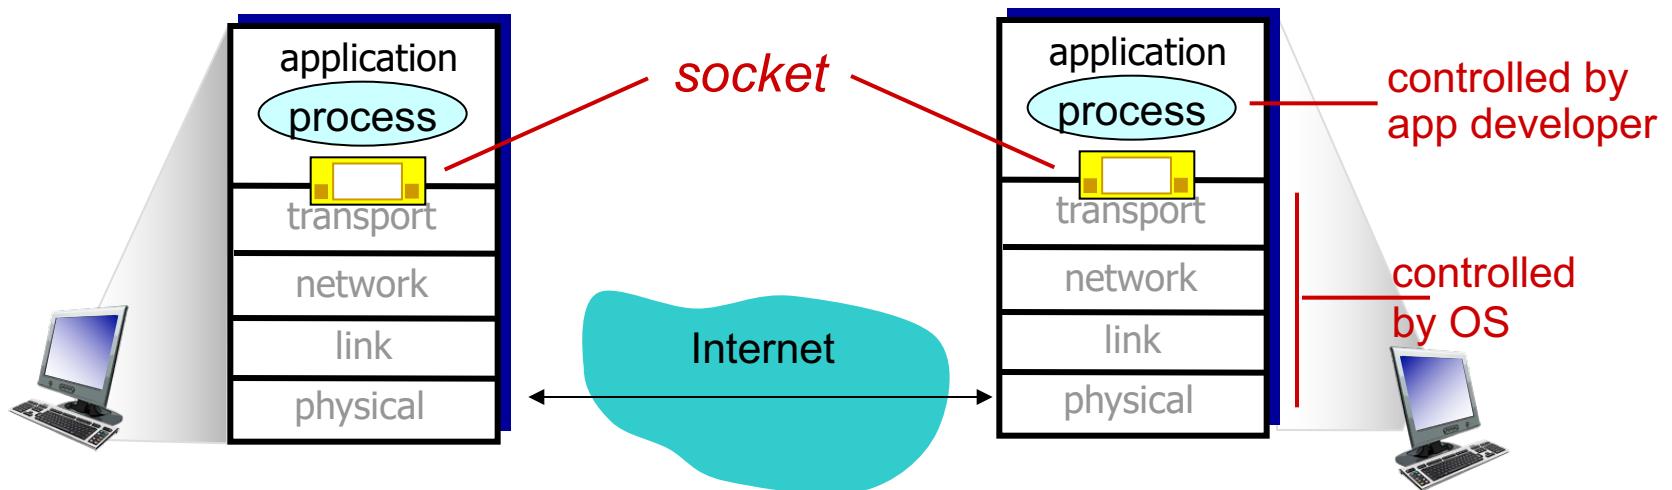
clients, servers

*client process*: process that initiates communication

*server process*: process that waits to be contacted

# Sockets

- ❖ Process sends messages into, and receives messages from, the network through a software interface called a **socket**.
- ❖ A process is like a house and its socket is like its door



# Addressing processes

- ❖ to receive messages, process must have *identifier*
- ❖ host device has unique 32-bit IP address
- ❖ *Q:* does IP address of host on which process runs suffice for identifying the process?
  - *A:* no, many processes can be running on same host
- ❖ *identifier* includes both IP address and port numbers associated with process on host.
- ❖ example port numbers:
  - HTTP server: 80
  - mail server: 25
- ❖ to send HTTP message to gaia.cs.umass.edu web server:
  - IP address: 128.119.245.12
  - port number: 80

# App-layer protocol defines

- ❖ types of messages exchanged,
  - e.g., request, response
- ❖ message syntax:
  - what fields in messages & how fields are delineated
- ❖ message semantics
  - meaning of information in fields
- ❖ rules for when and how processes send & respond to messages

**open protocols:**

- ❖ defined in RFCs
- ❖ allows for interoperability
- ❖ e.g., HTTP, SMTP

**proprietary protocols:**

- ❖ e.g., Skype

# What transport service does an app need?

## data integrity

- ❖ some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- ❖ other apps (e.g., audio) can tolerate some loss

## timing

- ❖ some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## throughput

- ❖ some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- ❖ other apps (“elastic apps”) make use of whatever throughput they get

## security

- ❖ encryption, data integrity,

...

# Transport service requirements: common apps

<b>application</b>	<b>data loss</b>	<b>throughput</b>	<b>time sensitive</b>
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100's msec
text messaging	no loss	elastic	yes and no

# Internet transport protocols services

## TCP service:

- ❖ *reliable transport* between sending and receiving process
- ❖ *flow control*: sender won't overwhelm receiver
- ❖ *congestion control*: throttle sender when network overloaded
- ❖ *does not provide*: timing, minimum throughput guarantee, security
- ❖ *connection-oriented*: setup required between client and server processes

## UDP service:

- ❖ *unreliable data transfer* between sending and receiving process
- ❖ *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

# Internet apps: application, transport protocols

	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

# Securing TCP

## TCP & UDP

- ❖ no encryption
- ❖ Clear text passwords sent into socket traverse Internet in clear text

## SSL

- ❖ provides encrypted TCP connection
- ❖ data integrity
- ❖ end-point authentication

## SSL is at app layer

- ❖ Apps use SSL libraries, which “talk” to TCP

## SSL socket API

- ❖ Clear text passwords sent into socket traverse Internet encrypted

# Chapter 2: outline

## 2.1 principles of network applications

- app architectures
- app requirements

## 2.2 Web and HTTP

## 2.3 FTP

## 2.4 electronic mail

- SMTP, POP3, IMAP

## 2.5 DNS

## 2.6 P2P applications

# Web and HTTP

- ❖ *web page* = Group of *objects*
- ❖ object can be HTML file, JPEG image, Java applet, audio file,...
- ❖ web page consists of *base HTML-file* which includes *several referenced objects*
- ❖ each object is addressable by a *URL*, e.g.,

www.someschool.edu/someDept/pic.gif

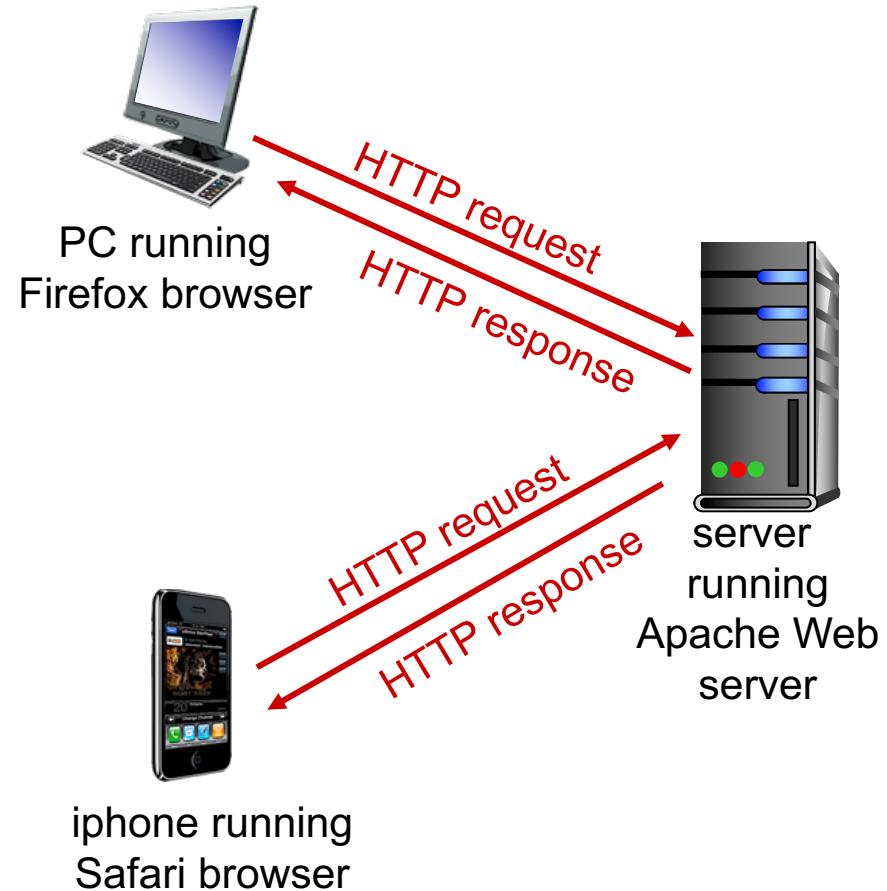
host name

path name

# HTTP overview

## HTTP: hypertext transfer protocol

- ❖ Web's application layer protocol
- ❖ client/server model
  - **client:** browser that requests, receives, (using HTTP protocol) and "displays" Web objects
  - **server:** Web server sends (using HTTP protocol) objects in response to requests



# HTTP overview (continued)

*uses TCP:*

1. client initiates TCP connection (creates socket) to server, port 80
2. server accepts TCP connection from client
3. HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
4. TCP connection closed

*HTTP is “stateless”*

- ❖ server maintains no information about past client requests

# HTTP connections

## *non-persistent HTTP*

- ❖ Open one TCP connection, Get one object, close
- ❖ downloading multiple objects required multiple connections

## *persistent HTTP*

- ❖ multiple objects can be sent over single TCP connection between client, server

# Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

Ia. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

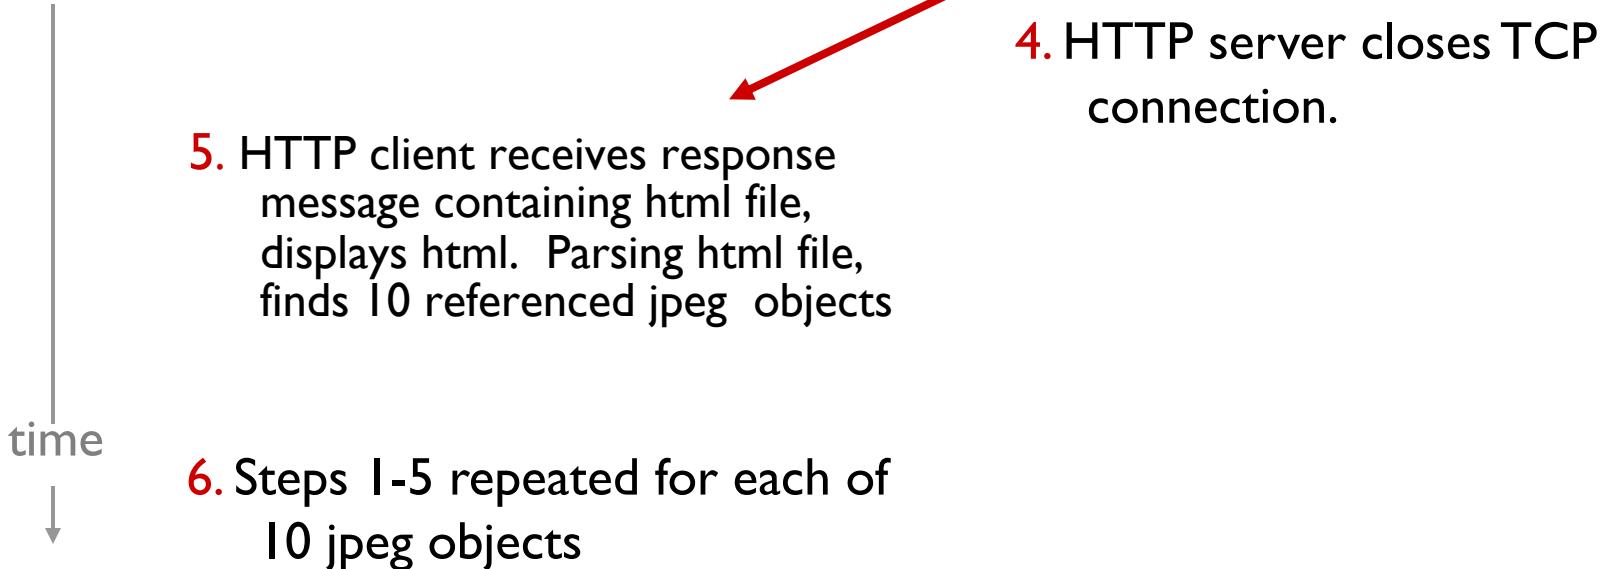
Ib. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time  
↓

# Non-persistent HTTP (cont.)



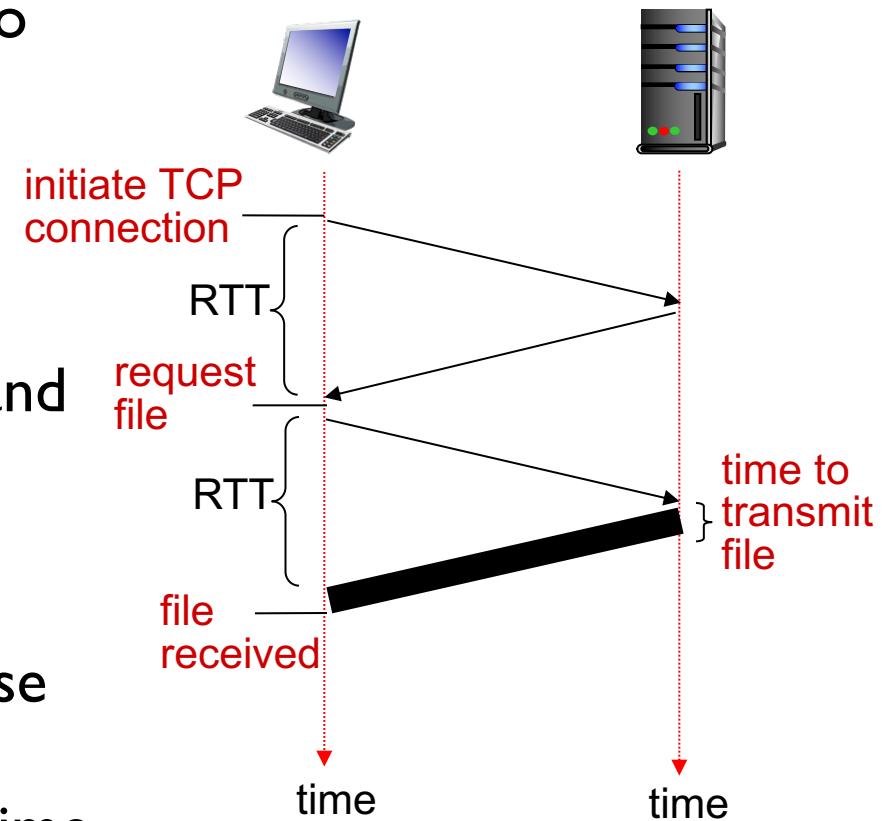
# Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time:

- ❖ one RTT to initiate TCP connection
- ❖ one RTT for HTTP request and first few bytes of HTTP response to return
- ❖ file transmission time
- ❖ non-persistent HTTP response time =

$$2\text{RTT} + \text{file transmission time}$$



# Persistent HTTP

## *non-persistent HTTP issues:*

- ❖ requires 2 RTTs per object
- ❖ OS overhead for each TCP connection
- ❖ browsers often open parallel TCP connections to fetch referenced objects

## *persistent HTTP:*

- ❖ server leaves connection open after sending response
- ❖ subsequent HTTP messages between same client/server sent over open connection
- ❖ client sends requests as soon as it encounters a referenced object
- ❖ as little as one RTT for all the referenced objects

# HTTP request message

- ❖ GET /somedir/page.html HTTP/1.1
- ❖ Host: www.someschool.edu
- ❖ Connection: close
- ❖ User-agent: Mozilla/5.0
- ❖ Accept-language: fr
- ❖ The request line has three fields:
  - The method field,
  - the URL field,
  - and the HTTP version field.

Request message

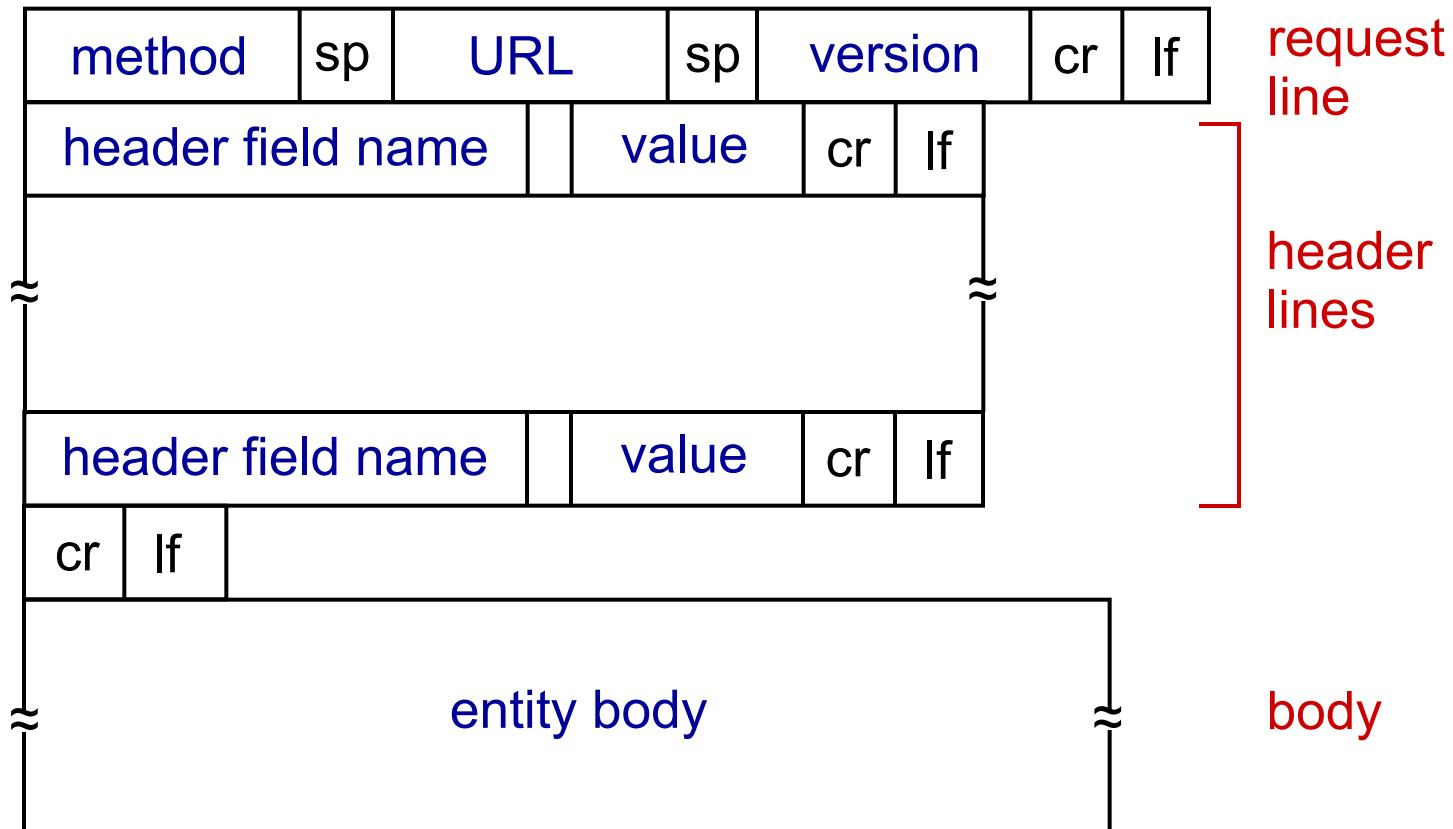
Header Lines



# Different Method Values

- ❖ **GET**: Get an object from server
- ❖ **POST**: Client often uses the POST method when the user fills out a form—for example, when a user provides search words to a search engine.
- ❖ **HEAD** :is similar to the GET. When a server receives a request with the HEAD method, it responds with an HTTP message but it leaves out the requested object. This method is used for debugging.
- ❖ **PUT**: is often used to upload an object to a specific path (directory) on a specific Webserver.
- ❖ **DELETE**: allows a user, or an application, to delete an object on a Web server.

# HTTP request message: general format



# HTTP response message

- ❖ HTTP/1.1 200 OK
- ❖ Connection: close
- ❖ Date: Tue, 09 Aug 2011 15:44:04 GMT
- ❖ Server: Apache/2.2.3 (CentOS)
- ❖ Last-Modified: Tue, 09 Aug 2011 15:11:03 GMT
- ❖ Content-Length: 6821
- ❖ Content-Type: text/html
- ❖ (data data data data data ...)

# HTTP response status codes

- ❖ status code appears in 1st line in server-to-client response message.
- ❖ some sample codes:

## **200 OK**

- request succeeded, requested object later in this msg

## **301 Moved Permanently**

- requested object moved, new location specified later in this msg  
(Location:)

## **400 Bad Request**

- request msg not understood by server

## **404 Not Found**

- requested document not found on this server

## **505 HTTP Version Not Supported**

# Trying out HTTP (client side) for yourself

## I. Telnet to your favorite Web server:

```
telnet gaia.cs.umass.edu 80
```

} opens TCP connection to port 80  
(default HTTP server port)  
at gaia.cs.umass.edu.  
anything typed in will be sent  
to port 80 at gaia.cs.umass.edu

## 2. type in a GET HTTP request:

```
GET /kurose_ross/interactive/index.php HTTP/1.1  
Host: gaia.cs.umass.edu
```

} by typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

## 3. look at response message sent by HTTP server! (or use Wireshark to look at captured HTTP request/response)

# User-server state: cookies

- ❖ many Web sites use cookies to keep track of users. Most major commercial Web sites use cookies today.

*four components:*

- 1) cookie header line of HTTP response message
- 2) cookie header line in next HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

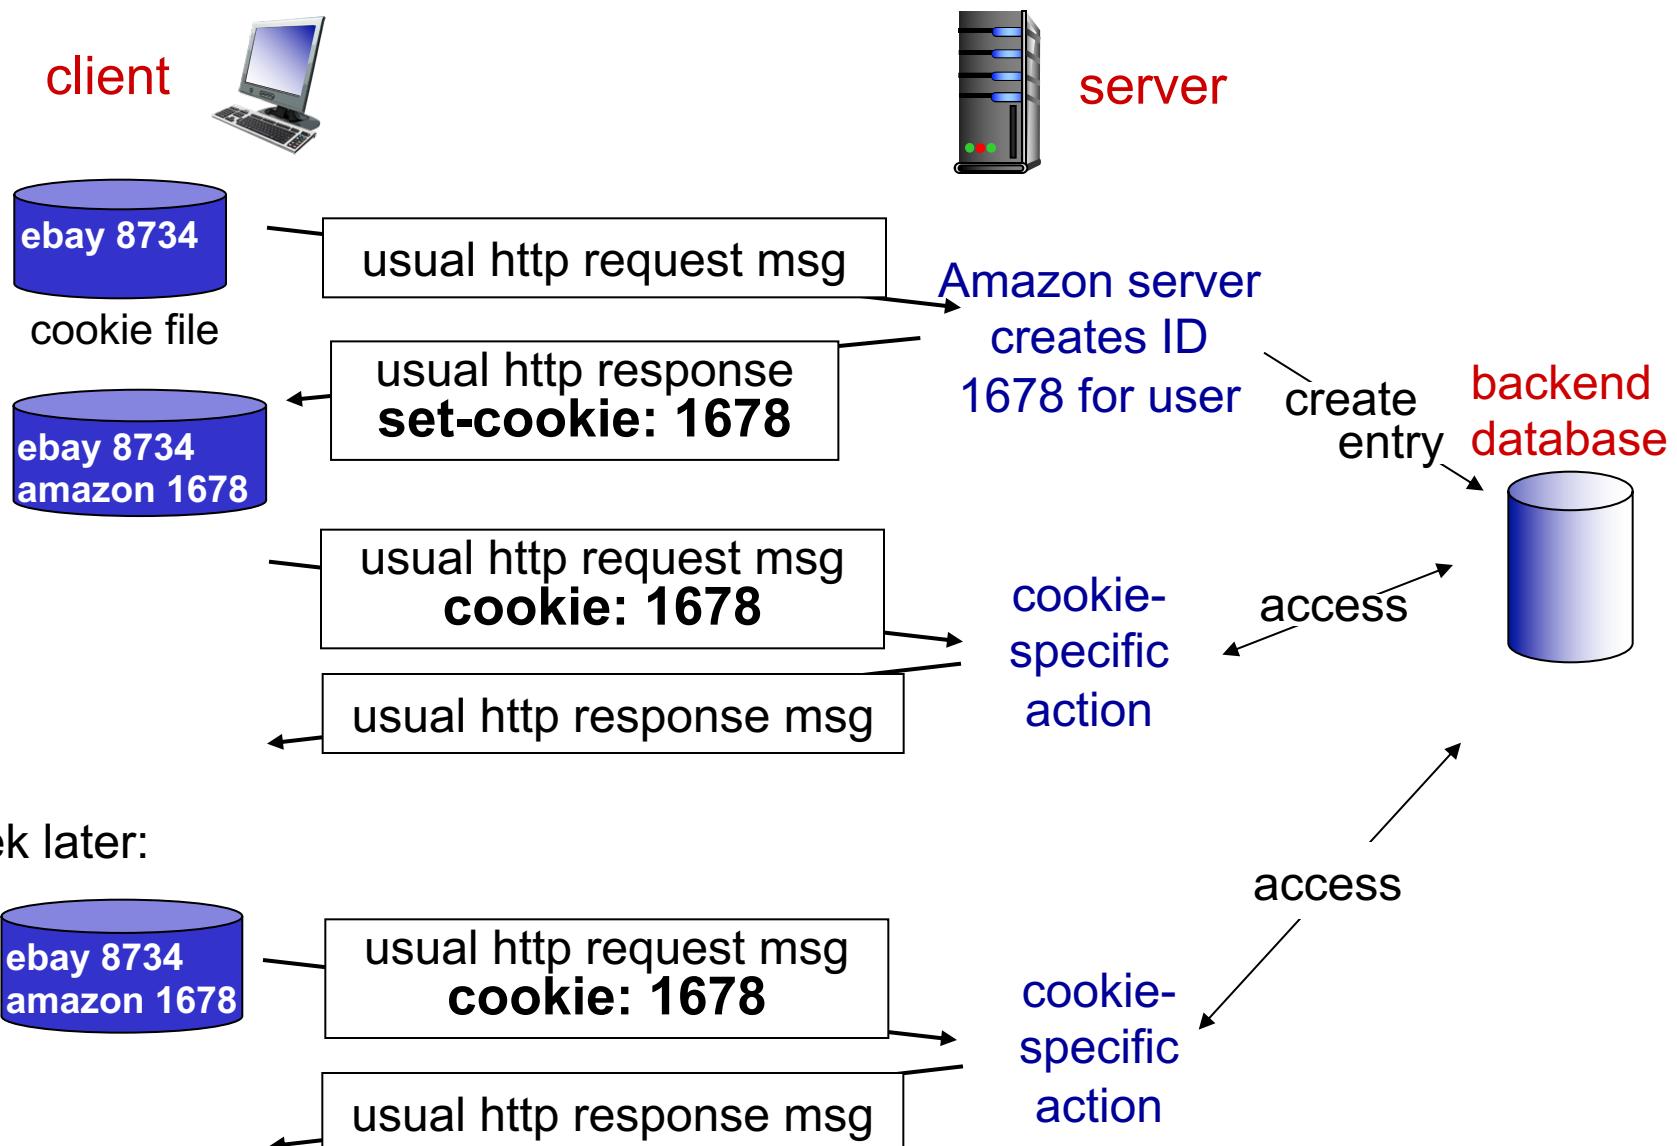
*example:*

- ❖ Susan always access Internet from PC
- ❖ visits specific e-commerce site for first time
- ❖ when initial HTTP requests arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

*what cookies can be used for:*

- ❖ authorization
- ❖ shopping carts
- ❖ recommendations

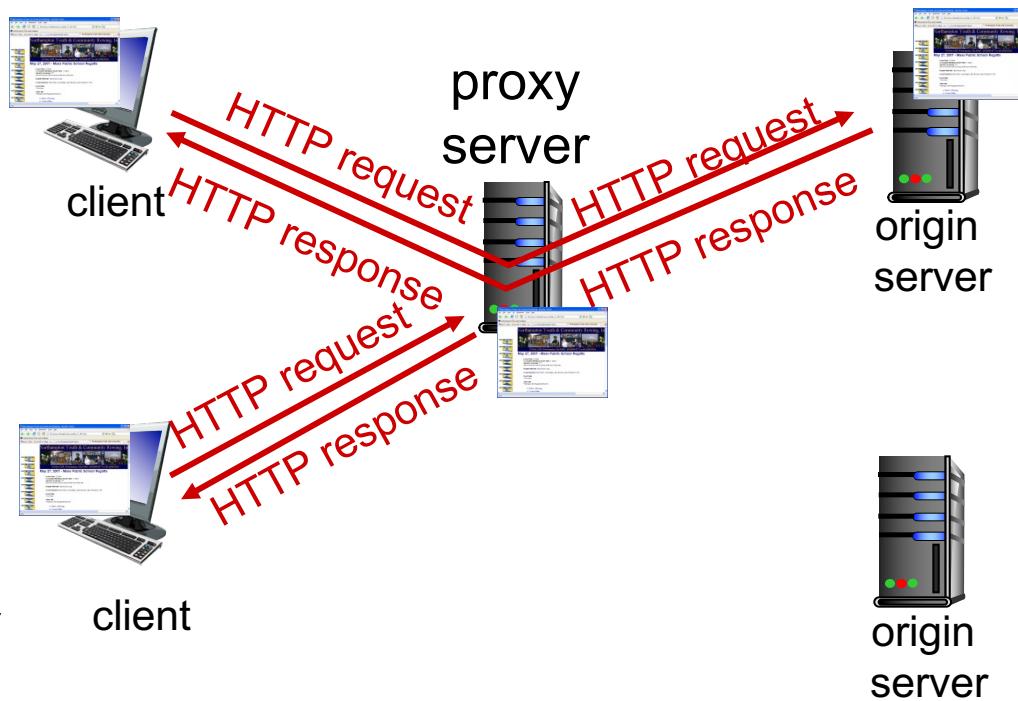
# Cookies: keeping “state” (cont.)



# Web caches (proxy server)

**goal:** satisfy client request without involving origin server

- ❖ All requests are sent to proxy server
- ❖ Proxy server caches objects
- ❖ Only new objects are requested from origin server
- ❖ Fast, Lower traffic on the link
- ❖ Typically a Web cache is purchased and installed by an ISP and a university



# Advantages of Web Caches

1. A Web cache can substantially reduce the response time for a client request, particularly if the bottleneck bandwidth between the client and the origin server
2. Web caches can substantially reduce traffic on access link to the Internet

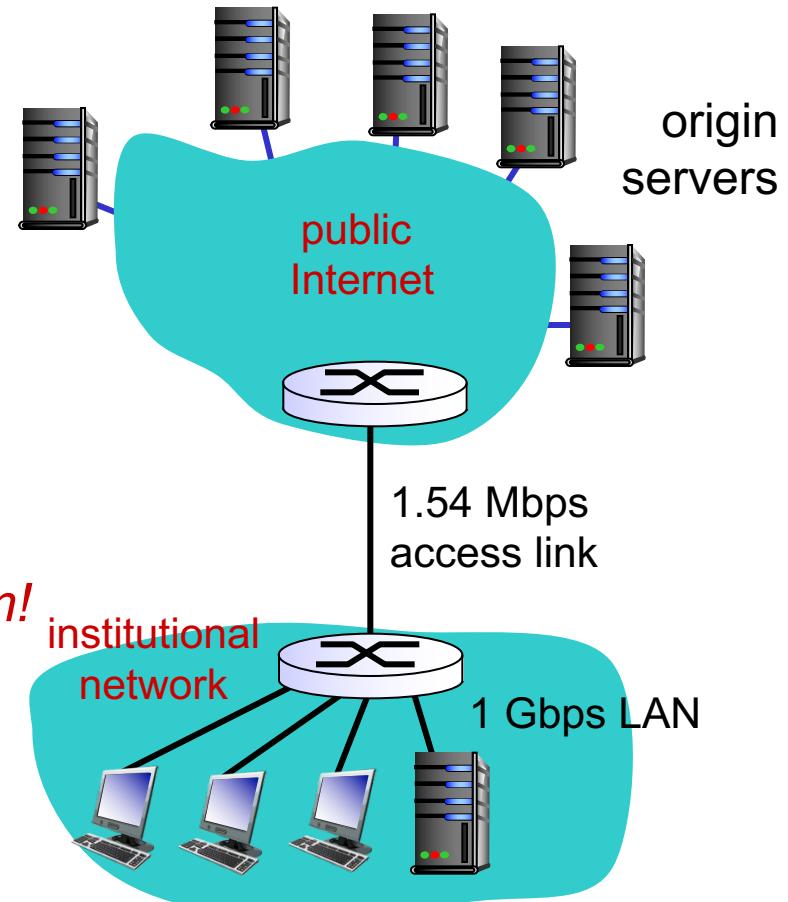
# Caching example:

## *assumptions:*

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

## *consequences:*

- LAN utilization: 15%
- access link utilization = **99%** *problem!*
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + usecs



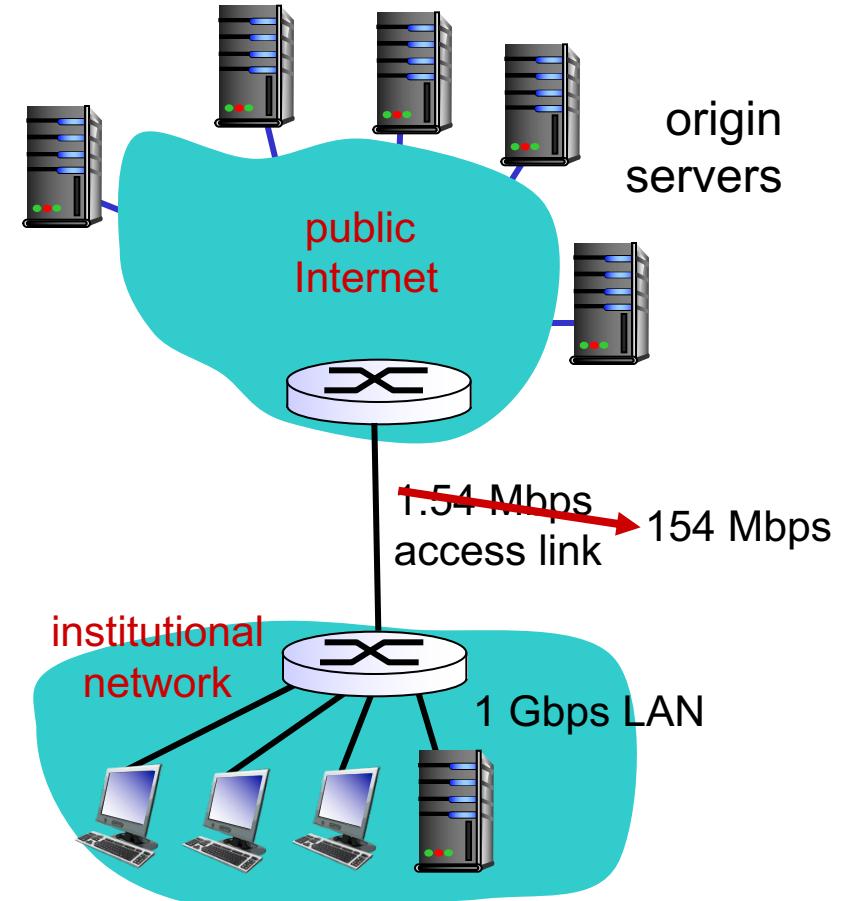
# Caching example: fatter access link

## *assumptions:*

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

## *consequences:*

- 154 Mbps
- LAN utilization: 15%
  - access link utilization = **99%**
  - total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + usecs
- 9.9%  
msecs



*Cost:* increased access link speed (not cheap!)

# Caching example: install local cache

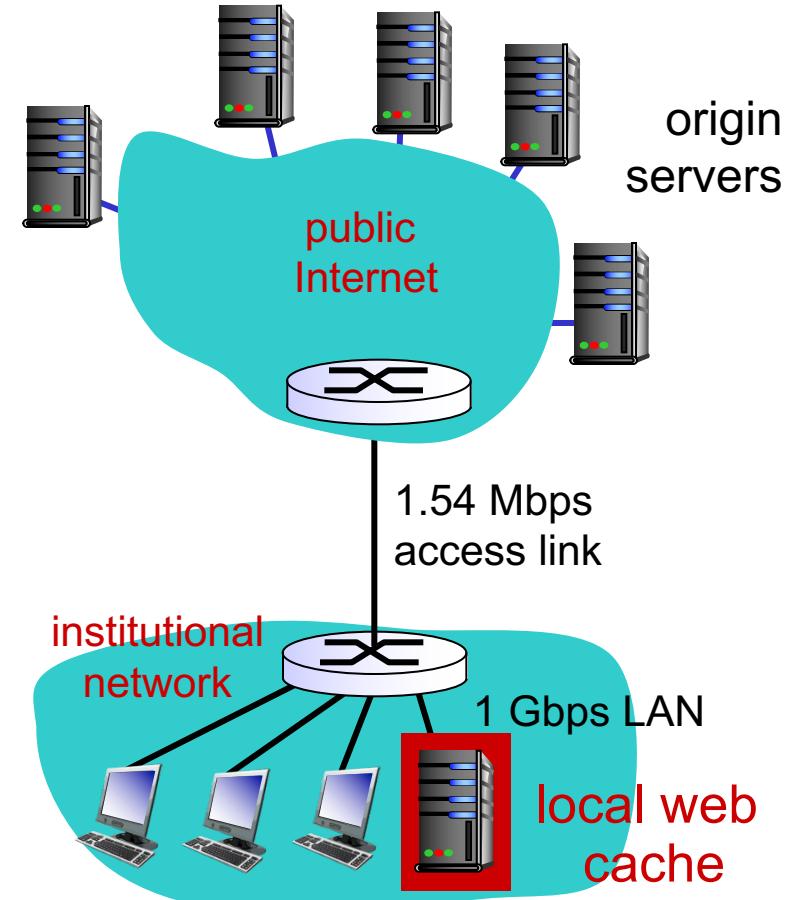
## *assumptions:*

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

## *consequences:*

- LAN utilization: 15%
- access link utilization = 100%
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + use<sup>~</sup>

*How to compute link utilization, delay?*

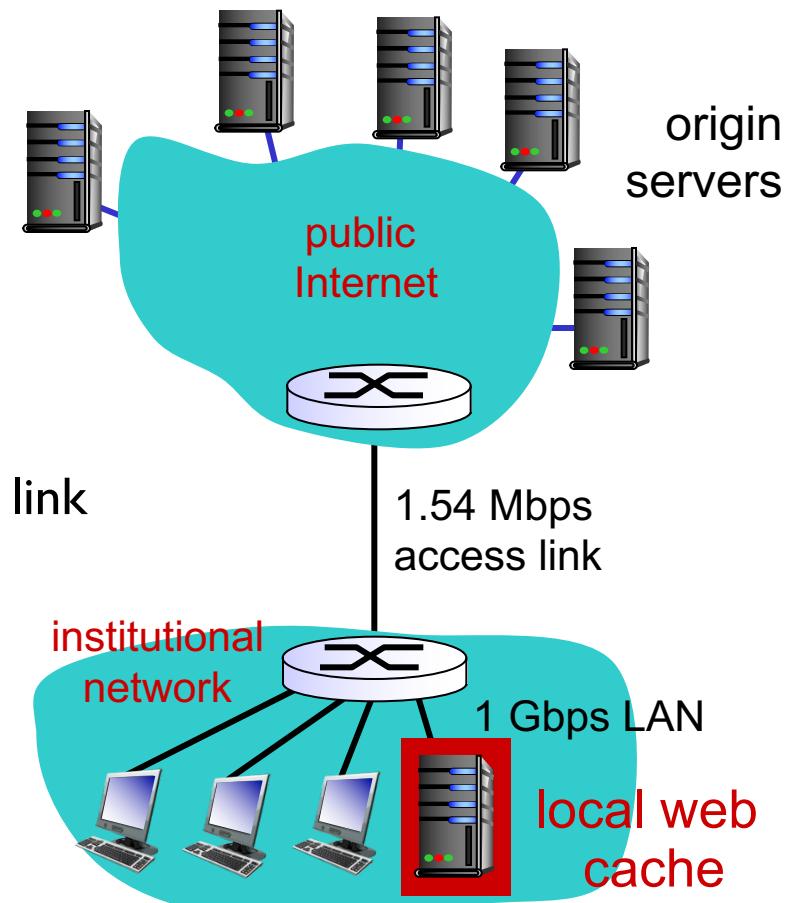


*Cost:* web cache (cheap!)

# Caching example: install local cache

*Calculating access link utilization, delay with cache:*

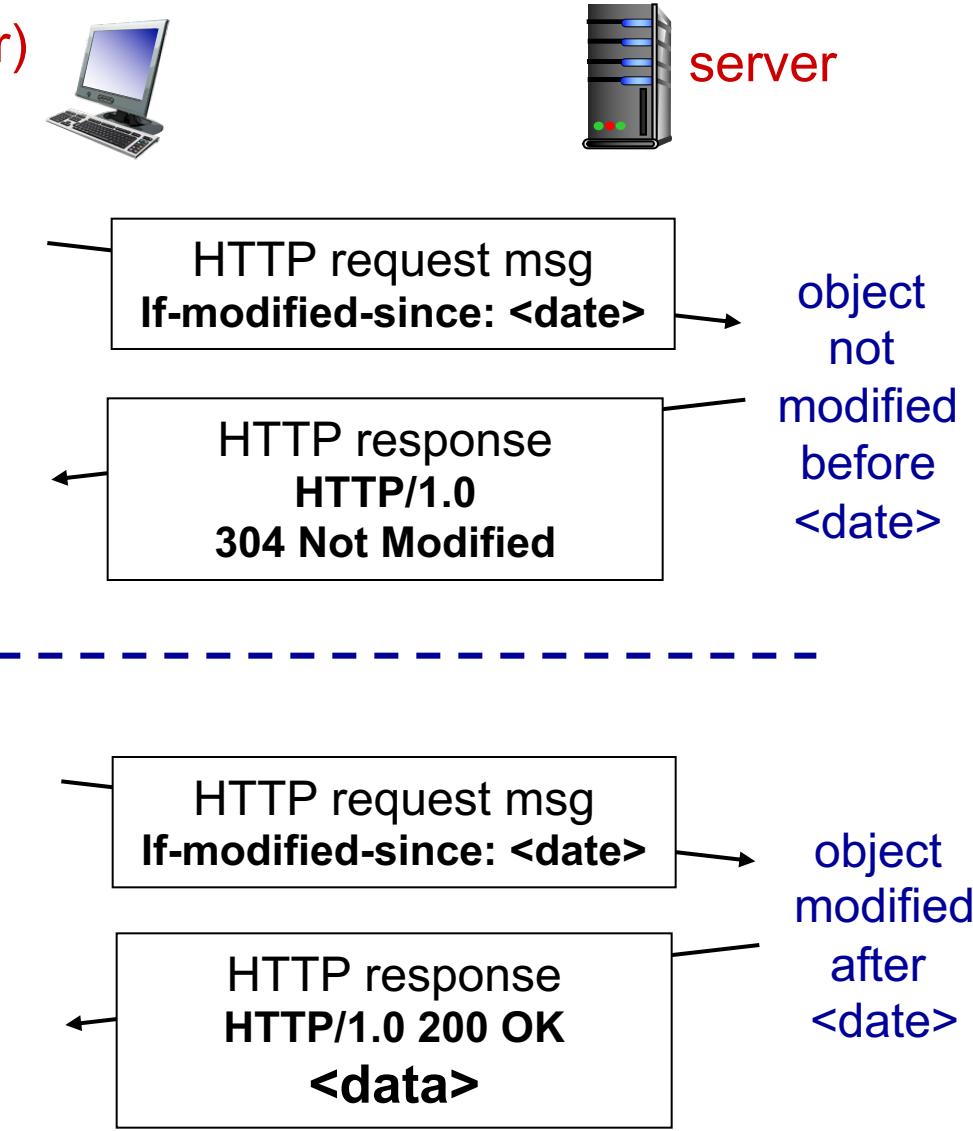
- ❖ suppose cache hit rate is 0.4
  - 40% requests satisfied at cache,  
60% requests satisfied at origin
- **access link utilization:**
  - 60% of requests use access link
- **data rate to browsers over access link**  
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$ 
  - utilization =  $0.9 / 1.54 = .58$
- **total delay**
  - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
  - less than with 154 Mbps link (and cheaper too!)



# Conditional GET

Client (cache Server)

- ❖ **Problem:** cache content may not be up-to-date
- ❖ **Goal:** don't send object if cache has up-to-date cached version
  - no object transmission delay
  - lower link utilization
- ❖ **cache:** specify date of cached copy in HTTP request  
**If-modified-since:**  
**<date>**
- ❖ **server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not**



# Chapter 2: outline

## 2.1 principles of network applications

- app architectures
- app requirements

## 2.2 Web and HTTP

## 2.3 FTP

## 2.4 electronic mail

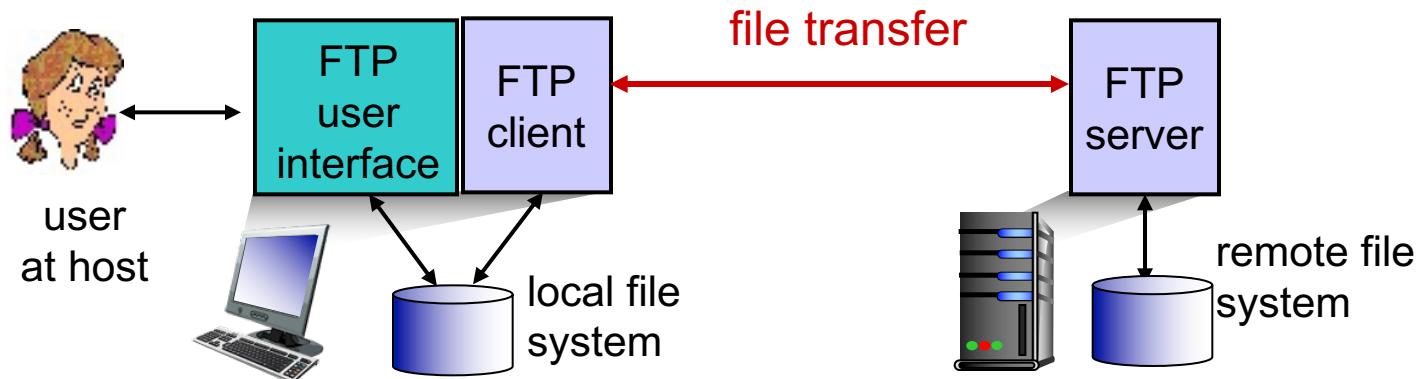
- SMTP, POP3, IMAP

## 2.5 DNS

## 2.6 P2P applications

## 2.7 socket programming with UDP and TCP

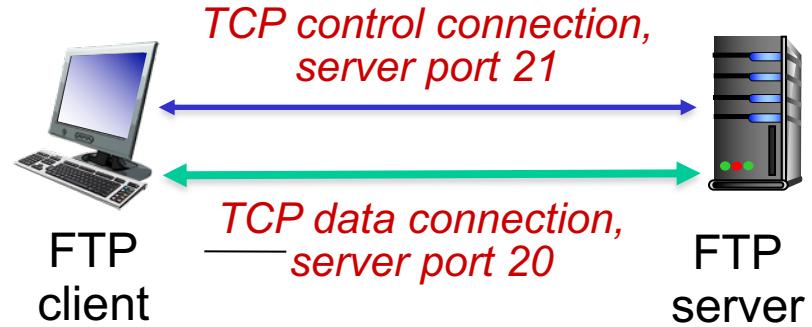
# FTP: the file transfer protocol



- ❖ transfer file to/from remote host
- ❖ client/server model
  - **client:** side that initiates transfer (either to/from remote)
  - **server:** remote host
- ❖ ftp: RFC 959
- ❖ ftp server: port 21

# FTP: separate control, data connections

- ❖ FTP client contacts FTP server at port 21, using TCP
- ❖ client authorized over control connection
- ❖ client browses remote directory, sends commands over control connection
- ❖ when server receives file transfer command, *server* opens 2<sup>nd</sup> TCP data connection (for file) to client
- ❖ after transferring one file, server closes data connection



- ❖ server opens another TCP data connection to transfer another file
- ❖ control connection: "*out of band*"
- ❖ FTP server maintains "state": current directory, earlier authentication

# FTP commands, responses

## *sample commands:*

- ❖ sent as ASCII text over control channel
- ❖ **USER** *username*
- ❖ **PASS** *password*
- ❖ **LIST** return list of file in current directory
- ❖ **RETR** *filename* retrieves (gets) file
- ❖ **STOR** *filename* stores (puts) file onto remote host

## *sample return codes*

- ❖ status code and phrase (as in HTTP)
- ❖ **331** Username OK, password required
- ❖ **125** data connection already open; transfer starting
- ❖ **425** Can't open data connection
- ❖ **452** Error writing file

# Chapter 2: outline

## 2.1 principles of network applications

- app architectures
- app requirements

## 2.2 Web and HTTP

## 2.3 FTP

## 2.4 electronic mail

- SMTP, POP3, IMAP

## 2.5 DNS

## 2.6 P2P applications

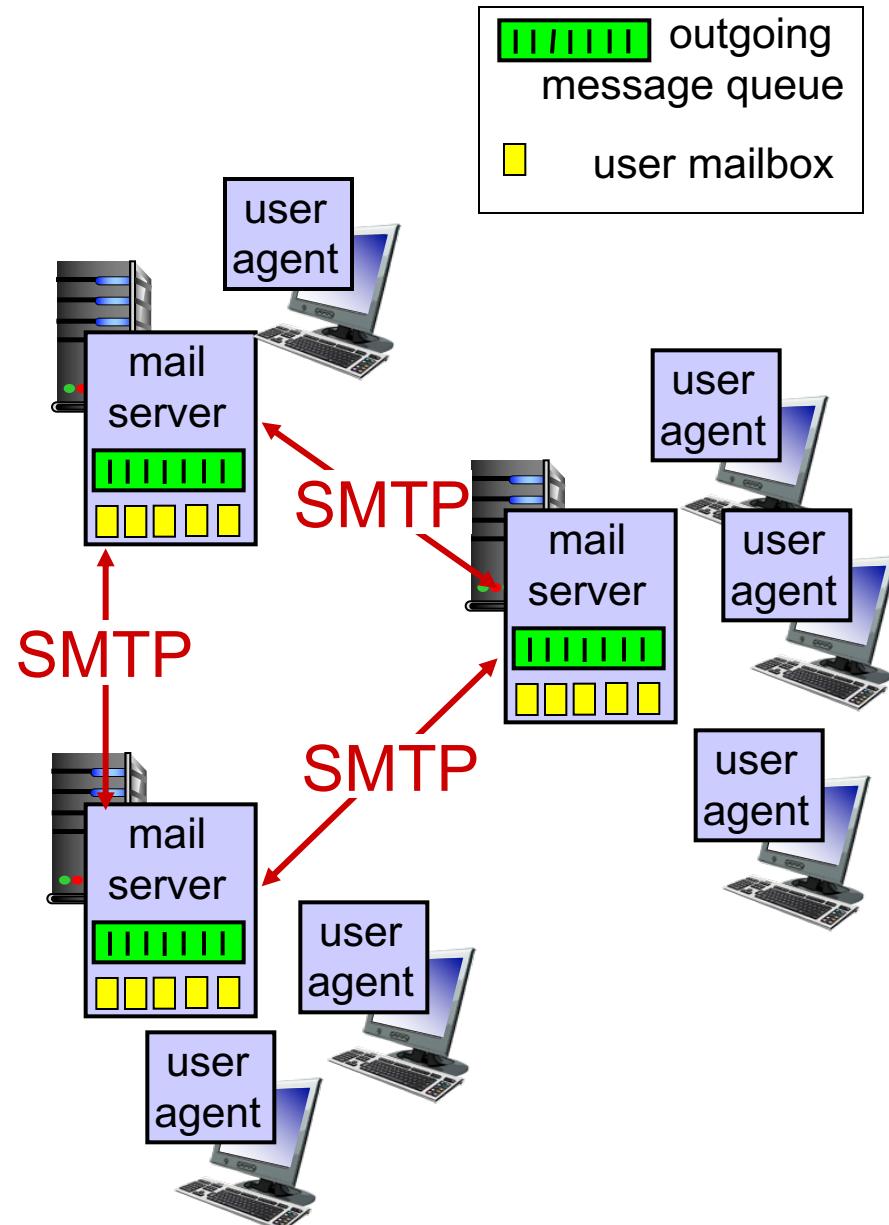
# Electronic mail

## *Three major components:*

1. user agents
2. mail servers
3. simple mail transfer protocol:  
SMTP

## *User Agent*

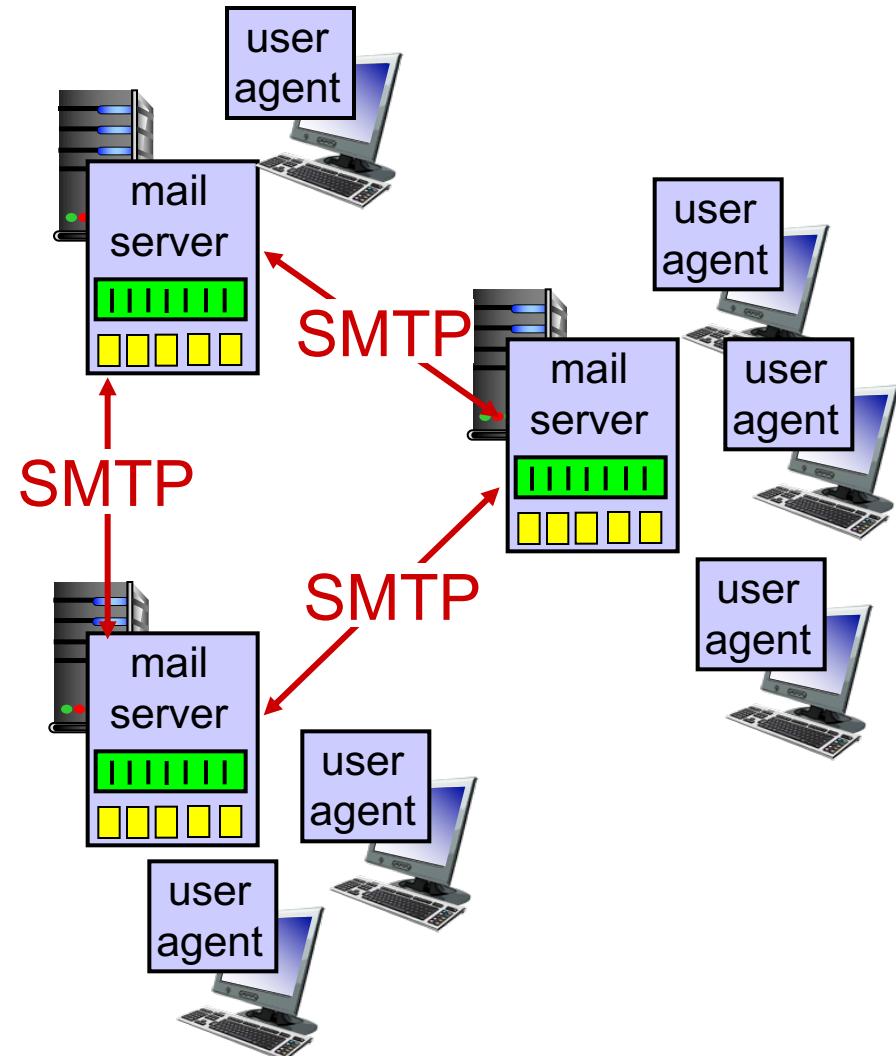
- ❖ a.k.a. “mail reader”
- ❖ composing, editing, reading mail messages
- ❖ e.g., Outlook, iPhone mail client
- ❖ outgoing, incoming messages stored on server



# Electronic mail: Mail servers

## Mail servers:

- ❖ **mailbox** contains incoming messages for user
- ❖ **message queue** of outgoing (to be sent) mail messages
- ❖ **SMTP protocol** between mail servers to send email messages
  - client: sending mail server
  - “server”: receiving mail server

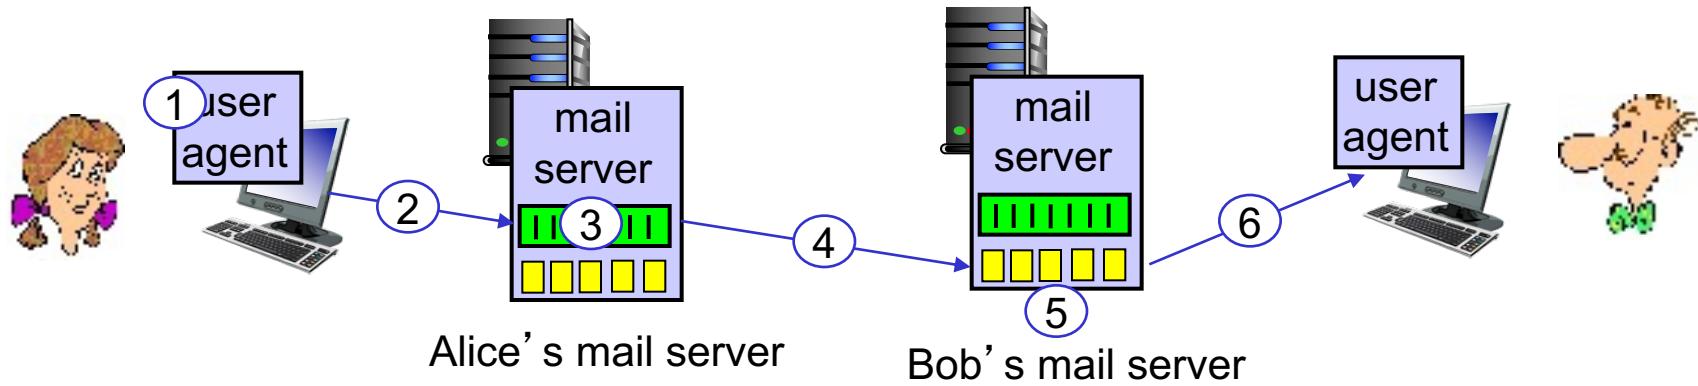


# Electronic Mail: SMTP [RFC 2821]

- ❖ uses TCP to reliably transfer email message from client to server, port 25
- ❖ three phases of transfer
  1. handshaking (greeting)
  2. transfer of messages
  3. closure
- ❖ command/response interaction (like HTTP, FTP)
  - commands: ASCII text
  - response: status code and phrase
- ❖ messages must be in 7-bit ASCII

# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



# Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

## Try SMTP interaction for yourself:

- ❖ **telnet servername 25**
- ❖ see 220 reply from server
- ❖ enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

# Comparing SMTP and HTTP

HTTP	SMTP
Persistent/Non-persistent	Persistent
Pull	Push
Accepts Binary objects	Accepts ASCII
One Object/Response	Multiple Objects/Messages

# Mail message format

SMTP: protocol for  
exchanging email msgs

RFC 822: standard for text  
message format:

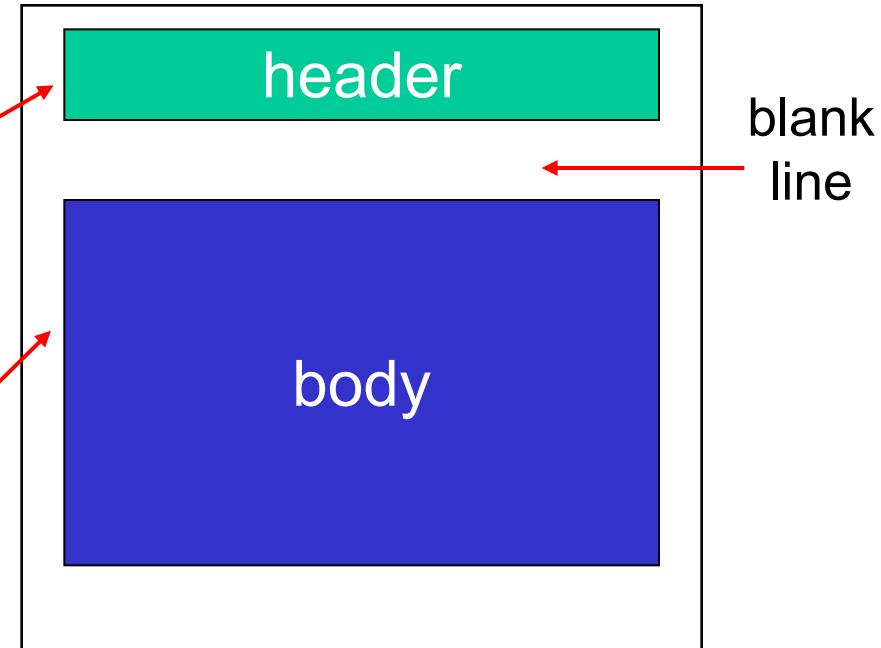
❖ header lines, e.g.,

- To:
- From:
- Subject:

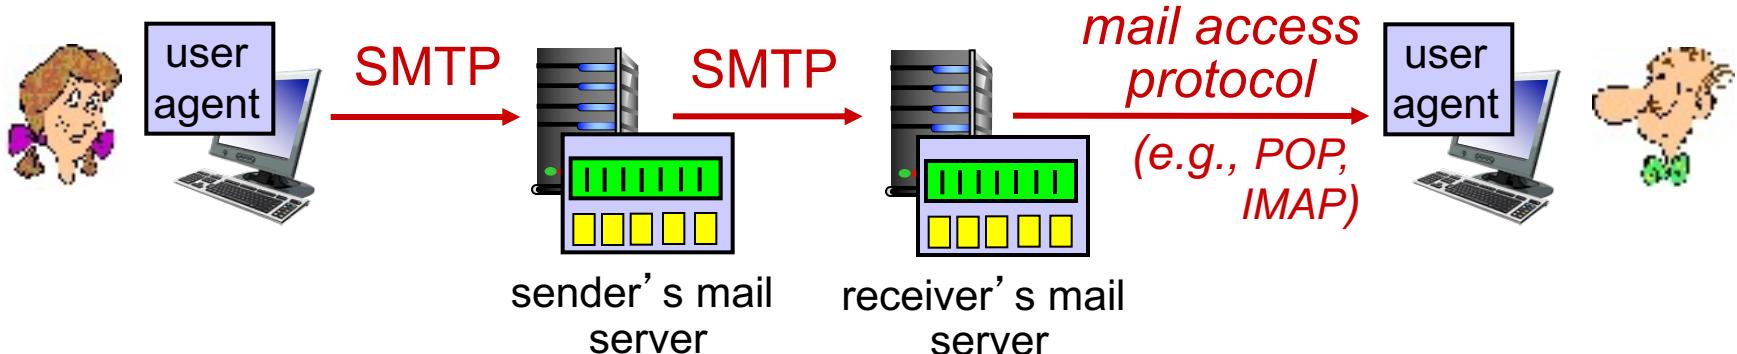
*different from SMTP MAIL  
FROM, RCPT TO:  
commands!*

❖ Body: the “message”

- ASCII characters only



# Mail access protocols



- ❖ **SMTP:** delivery/storage to receiver's server
- ❖ mail access protocol: retrieval from server
  - **POP:** Post Office Protocol [RFC 1939]: authorization, download
  - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored msgs on server
  - **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.

# POP3 protocol

## *authorization phase*

- ❖ client commands:
  - **user**: declare username
  - **pass**: password
- ❖ server responses
  - **+OK**
  - **-ERR**

## *transaction phase, client:*

- ❖ **list**: list message numbers
- ❖ **retr**: retrieve message by number
- ❖ **dele**: delete
- ❖ **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

# POP3 (more) and IMAP

## *more about POP3*

- ❖ previous example uses POP3 “download and delete” mode
  - Bob cannot re-read e-mail if he changes client
- ❖ POP3 “download-and-keep”: copies of messages on different clients
- ❖ POP3 is stateless across sessions

## *IMAP*

- ❖ keeps all messages in one place: at server
- ❖ allows user to organize messages in folders
- ❖ keeps user state across sessions:
  - names of folders and mappings between message IDs and folder name

# Chapter 2: outline

## 2.1 principles of network applications

- app architectures
- app requirements

## 2.2 Web and HTTP

## 2.3 FTP

## 2.4 electronic mail

- SMTP, POP3, IMAP

## 2.5 DNS

## 2.6 P2P applications

# DNS: Domain Name System

- ❖ There are two ways to identify a host
  1. by a hostname
    - “name”, e.g.,  
www.yahoo.com - used by humans
  2. and by an IP address
    - IP address (32 bit)
    - Example: 121.7.106.83, where each period separates one of the bytes expressed in decimal notation from 0 to 255.
    - “Q: how to map between IP address and name, and vice versa ?

DNS servers translate a host name to IP address

## *Domain Name System:*

- ❖ *Is a distributed database* implemented in hierarchy of many *name servers*
- ❖ *application-layer protocol* use *DNS* : to *resolve names* (address/name translation)

# DNS: services, structure

## *DNS services*

- ❖ hostname to IP address translation
- ❖ host aliasing
  - alias names: . host with a complicated hostname can have one or more alias names.
  - Example: relay1.west-coast.enterprise.com could have, aliase such as enterprise.com
- ❖ mail server aliasing to have short email address
- ❖ load distribution
  - replicated Web servers: many IP addresses correspond to one name. Rotated each time.

# Load Distribution

```
b: Non-authoritative answer:  
Name: google.com  
 Addresses: 2607:f8b0:4002:c07::64  
           74.125.196.138  
           74.125.196.100  
           74.125.196.101  
           74.125.196.102  
           74.125.196.139  
           74.125.196.113  
  
C:\Users\keshtgari>nslookup google.com  
Server: dns4.uga.edu  
Address: 128.192.1.19  
  
Non-authoritative answer:  
Name: google.com  
 Addresses: 2607:f8b0:4002:c07::64  
           74.125.196.101  
           74.125.196.102  
           74.125.196.100  
           74.125.196.138  
           74.125.196.113  
           74.125.196.139
```

# nslookup

```
Microsoft Windows [Version 10.0.10586]
(c) 2015 Microsoft Corporation. All rights reserved.
```

```
C:\Users\keshtgari>nslookup www.uga.edu
```

```
Server: dns4.uga.edu
Address: 128.192.1.19
```

```
Name: www.uga.edu
```

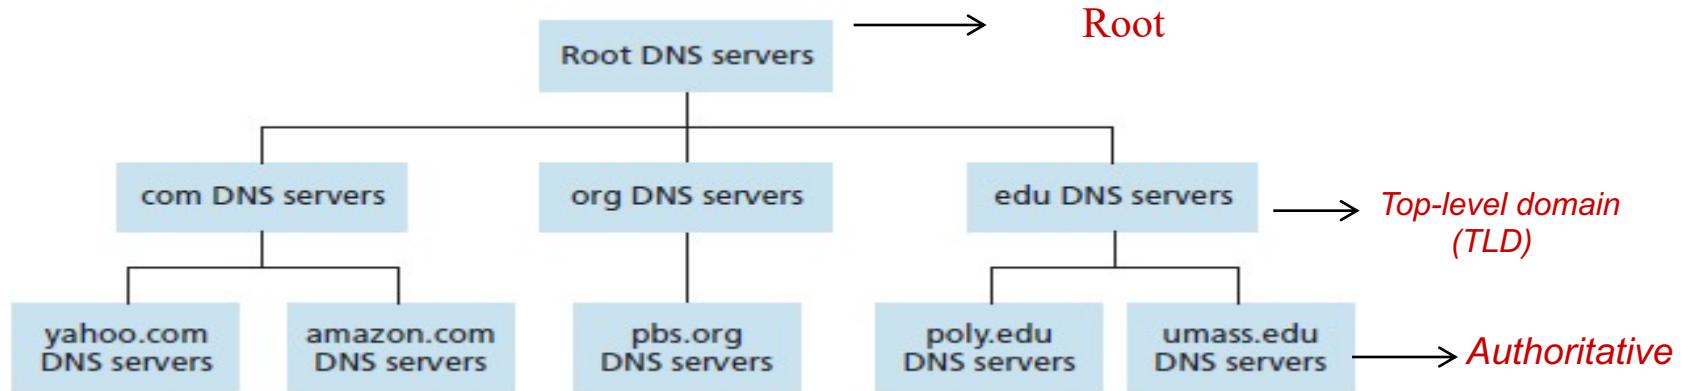
```
Address: 128.192.225.20
```

# DNS: services, structure

## *why not centralize DNS?*

- ❖ single point of failure: If the DNS server crashes, so does the entire Internet!
- ❖ traffic volume: HTTP requests and e-mail messages generated could be from hundreds of millions of hosts.
- ❖ distant centralized database:
- ❖ maintenance

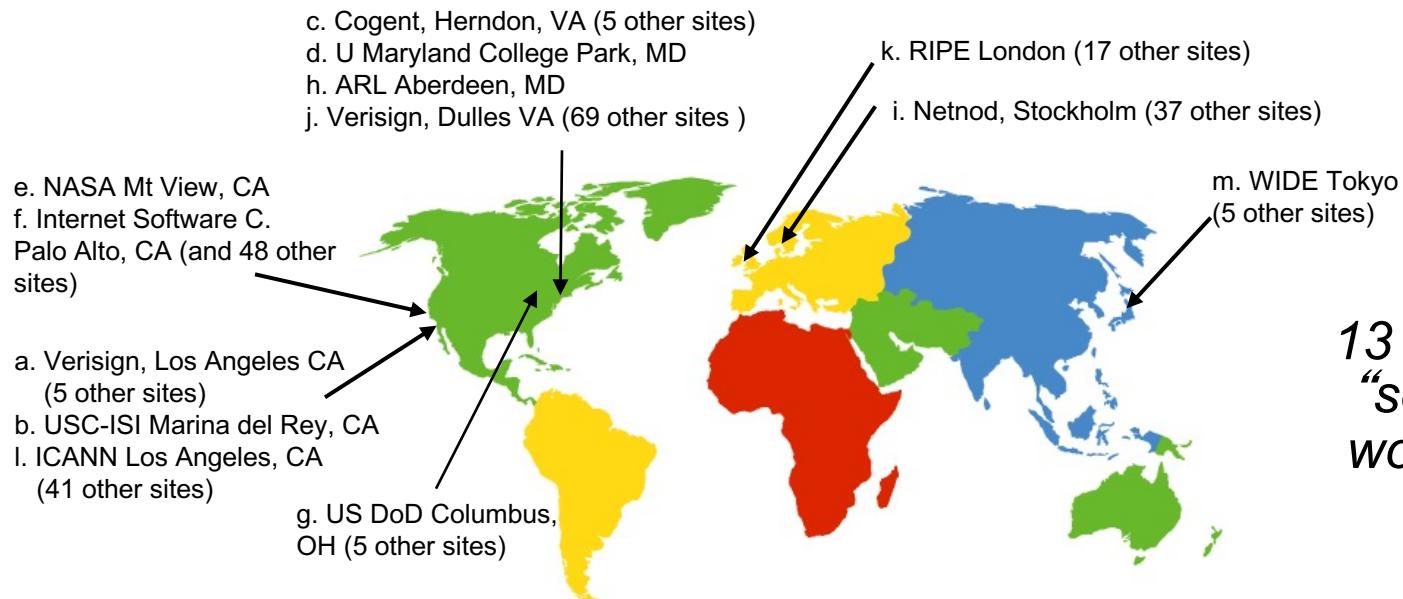
# DNS: a distributed, hierarchical database



*Root : 13 root DNS servers (labeled A through M), most of which are located in North America. Each “server” is actually a network of replicated servers, for both security and reliability purposes. All together, there are 247 root servers as of fall 2011.*

# DNS: Root name servers

- ❖ contacted by local name server that can not resolve name
- ❖ root name server:
  - contacts authoritative name server if name mapping not known
  - gets mapping
  - returns mapping to local name server



*13 root name  
“servers”  
worldwide*

# TLD, authoritative servers

## *Top-level domain (TLD) servers:*

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

## *Authoritative DNS servers:*

- ❖ Every organization with publicly accessible hosts (such as Web servers and mail servers) on the Internet must provide publicly accessible DNS records that map the names of those hosts to IP addresses.
  - can be maintained by organization or service provider

# Local DNS name server

- ❖ does not strictly belong to hierarchy
- ❖ each ISP (residential ISP, company, university) has one
  - also called “default name server”
- ❖ when host makes DNS query, query is sent to its local DNS server
  - has local cache of recent name-to-address translation pairs (but may be out of date!)
  - acts as proxy, forwards query into hierarchy

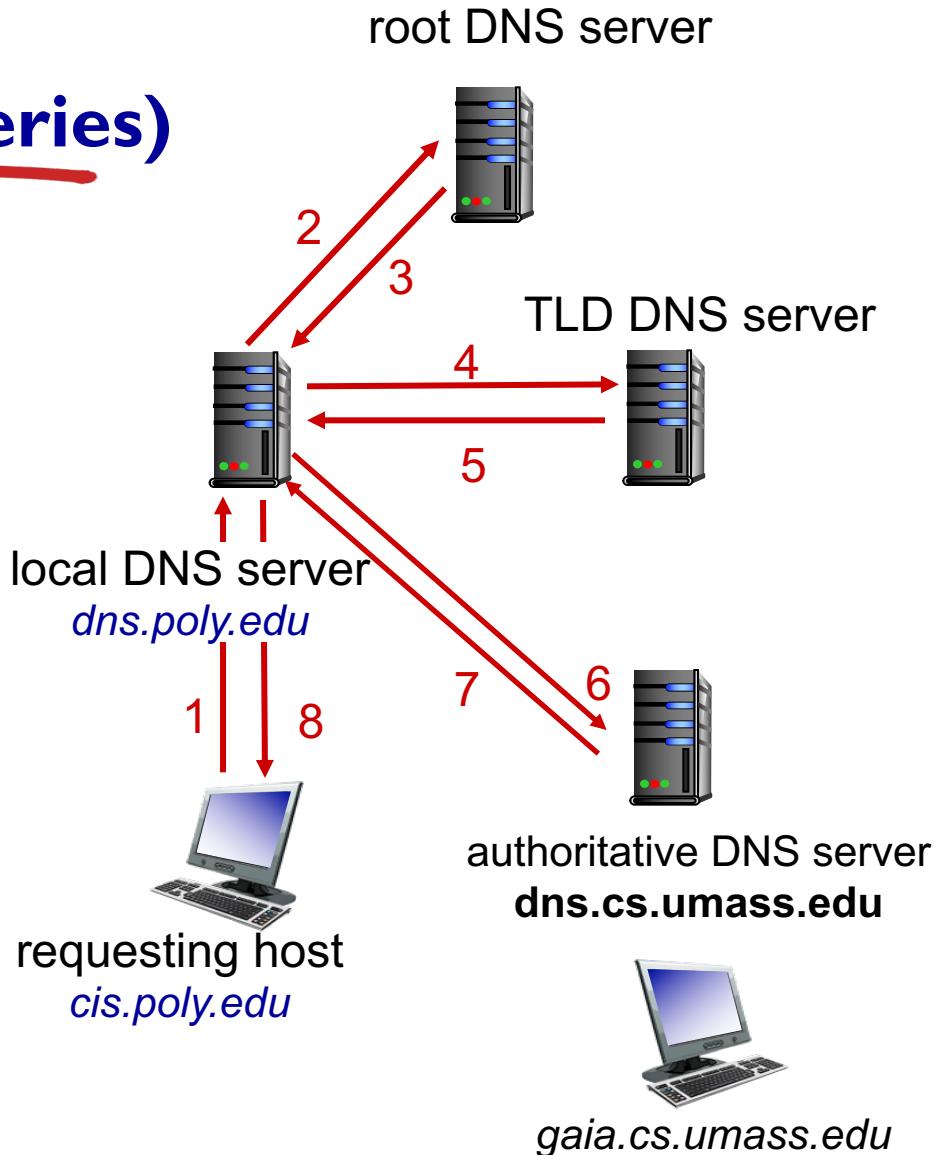
# DNS name resolution

## example(iterative queries)

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

### *iterated query:*

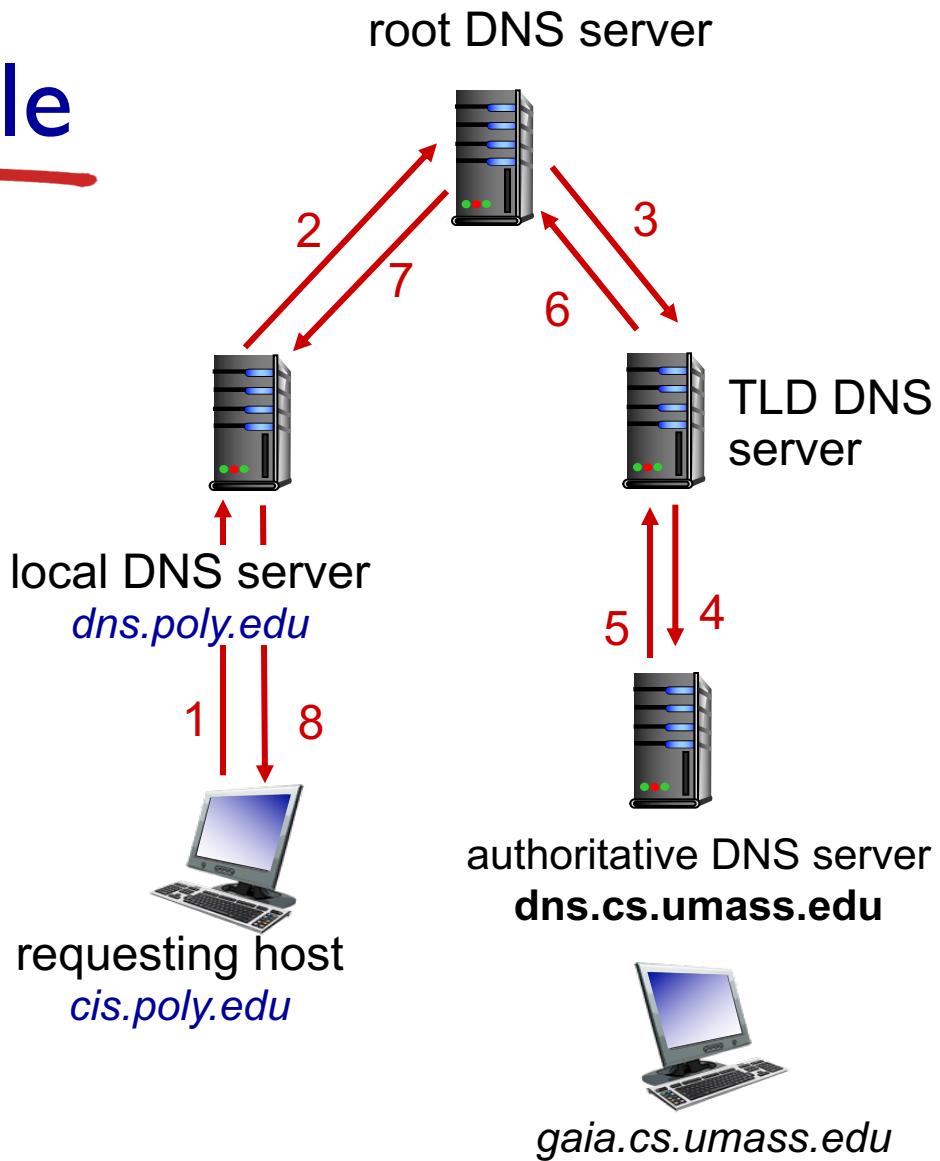
- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



# DNS name resolution example

*recursive query:*

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?



# DNS: caching, updating records

- ❖ once (any) name server learns mapping, it *caches* mapping
  - cache entries timeout (disappear) after some time (TTL)
  - TLD servers typically cached in local name servers
    - thus root name servers not often visited
- ❖ cached entries may be *out-of-date* (best effort name-to-address translation!)
  - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- ❖ update/notify mechanisms proposed IETF standard
  - RFC 2136

# DNS records

**DNS:** distributed db storing resource records (**RR**)

RR format: `(name, value, type, ttl)`

## type=A

- **name** is hostname
- **value** is IP address
- (relay1.bar.foo.com, 145.37.93.126, A)

## type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain
- (foo.com, dns.foo.com, NS)

## type=CNAME

- **name** is alias name for some “canonical” (the real) name
- (foo.com, relay1.bar.foo.com, CNAME)

## type=MX

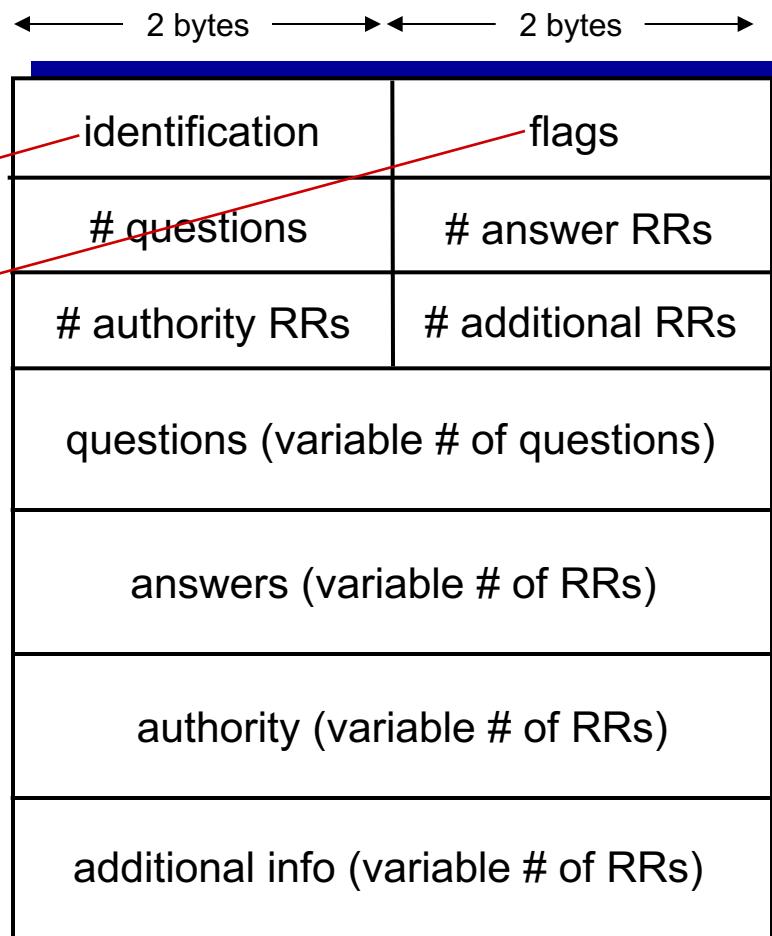
- **value** is name of mailserver associated with **name**
- (foo.com, mail.bar.foo.com, MX)

# DNS protocol, messages

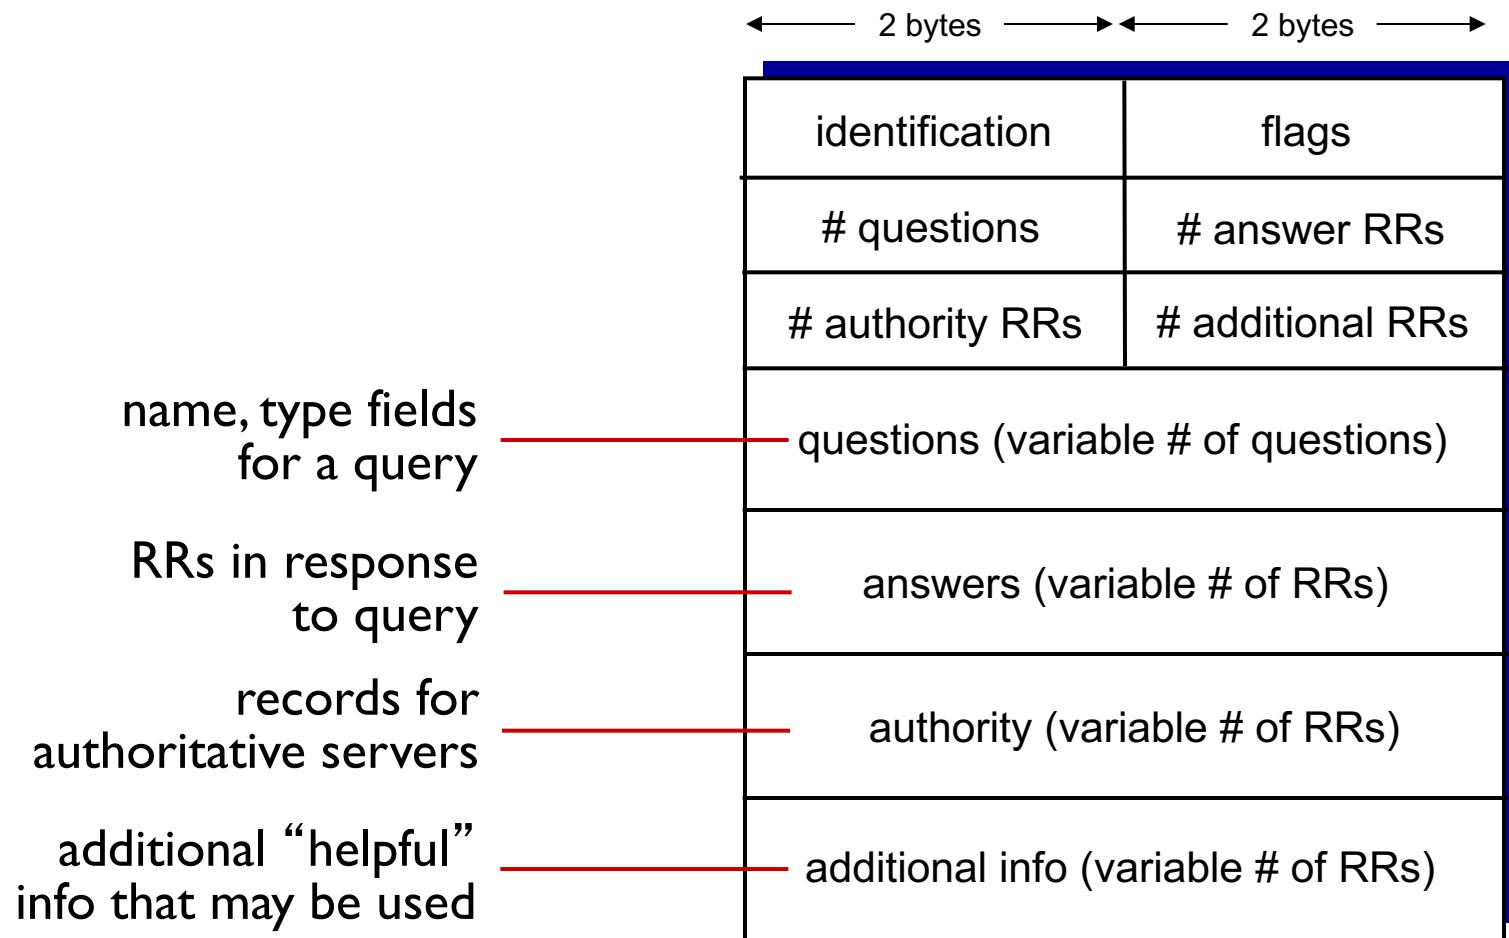
- ❖ *query* and *reply* messages, both with same *message format*

msg header

- ❖ **identification:** 16 bit # for query, reply to query uses same #
- ❖ **flags:**
  - query or reply
  - recursion desired
  - recursion available
  - reply is authoritative



# DNS protocol, messages



# Inserting records into DNS

## ~~(Example)~~

- ❖ example: new startup “Network Utopia”
- ❖ **First Step:** register domain name (`networkuptopia.com`) at **DNS registrar** (e.g., Network Solutions). A **registrar** is a commercial entity that verifies the uniqueness of the domain name, enters the domain name into the DNS database.
- ❖ **Second Step:** provide names, IP addresses of authoritative name server (primary and secondary) to the registrar. For example `dns1.networkutopia.com`
- ❖ registrar inserts two RRs into .com TLD server:  
`(networkutopia.com, dns1.networkutopia.com, NS)`  
`(dns1.networkutopia.com, 212.212.212.1, A)`
- ❖ You should also create authoritative server type A record for your webserver (`www.networkutopia.com`); Type MX resource record for your mail server (`mail.networkutopia.com`) are entered into your authoritative DNS servers

# Attacking DNS

## DDoS attacks

- ❖ Bombard root servers with traffic
  - Not successful to date
  - Traffic Filtering
  - Local DNS servers cache IPs of TLD servers, allowing root server bypass
- ❖ DNS cache poisoning – A server gives **wrong answer**

# Chapter 2: outline

## 2.1 principles of network applications

- app architectures
- app requirements

## 2.2 Web and HTTP

## 2.3 FTP

## 2.4 electronic mail

- SMTP, POP3, IMAP

## 2.5 DNS

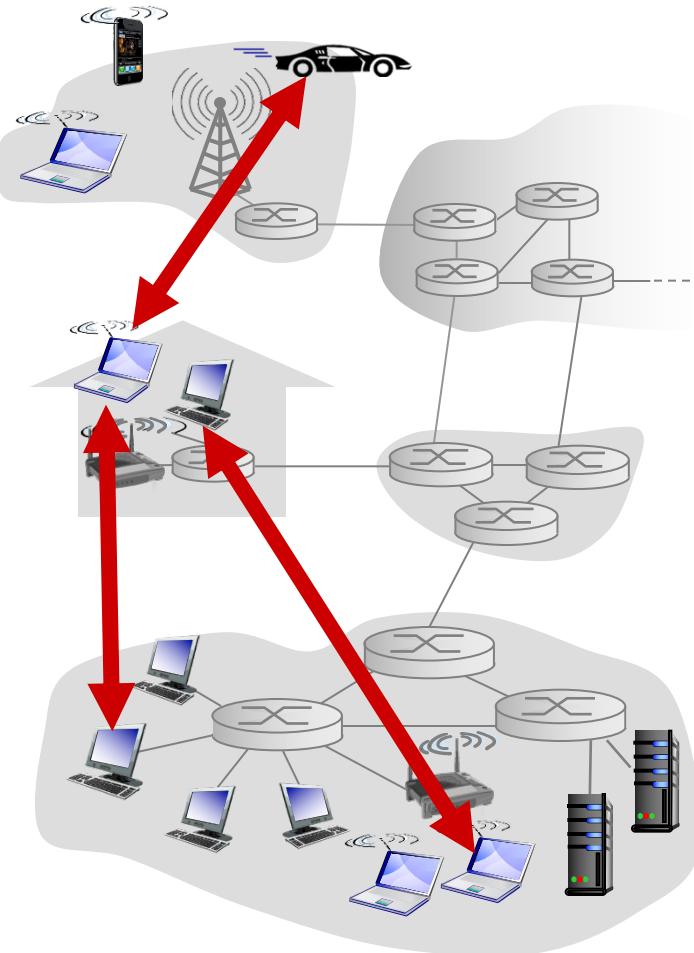
## 2.6 P2P applications

# Pure P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers are intermittently connected and change IP addresses

## *examples:*

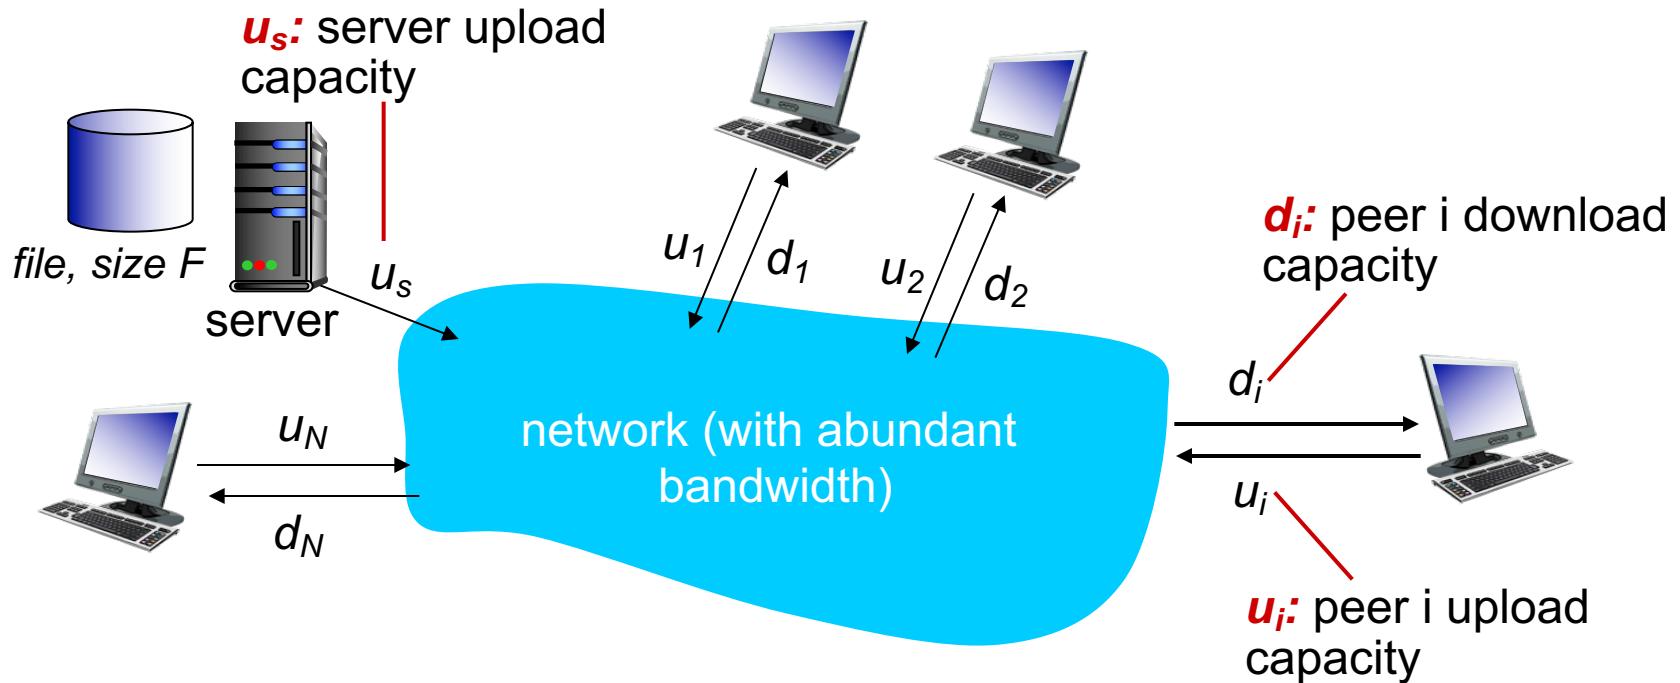
- file distribution (BitTorrent)
- VoIP (Skype)



# File distribution: client-server vs P2P

Question: how much time to distribute file (size  $F$ ) from one server to  $N$  peers?

- peer upload/download capacity is limited resource



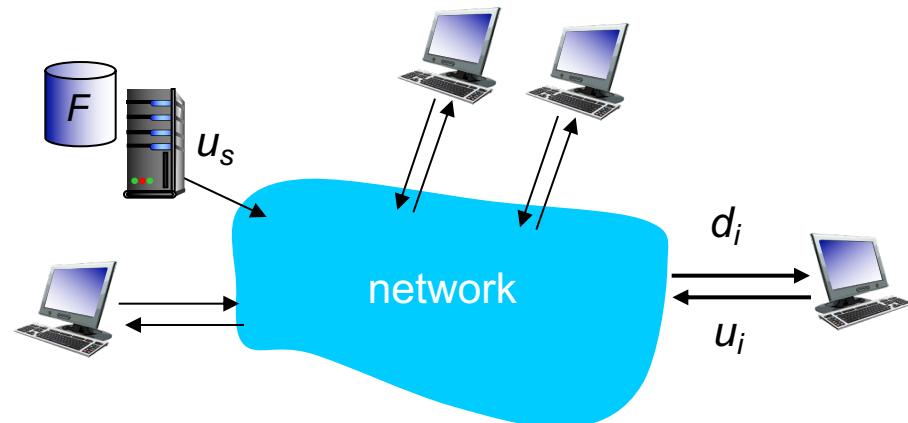
# File distribution time: client-server

- ❖ **server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  $NF/u_s$

- ❖ **client:** each client must download file copy

- $d_{\min}$  = min client download rate
- min client download time:  $F/d_{\min}$

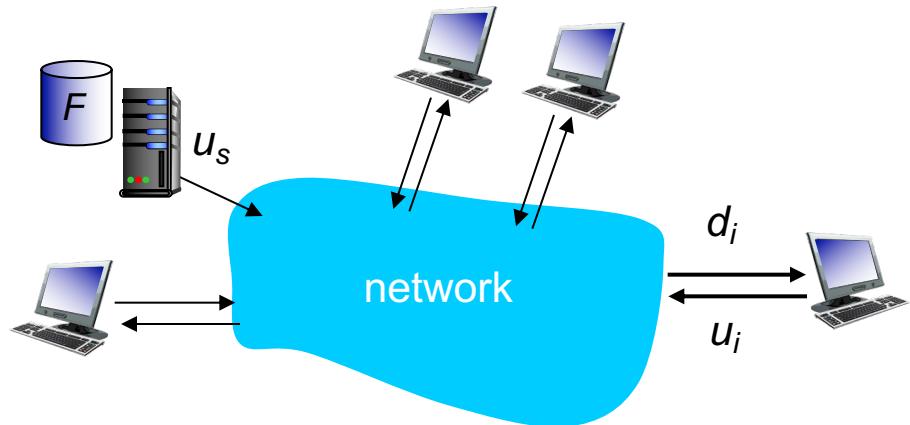


*time to distribute  $F$   
to  $N$  clients using  
client-server approach*  $D_{c-s} > \max\{NF/u_s, F/d_{\min}\}$

increases linearly in  $N$

# File distribution time: P2P

- ❖ **server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- ❖ **client:** each client must download file copy
  - min client download time:  $F/d_{\min}$



- ❖ **clients:** as aggregate must download  $NF$  bits
  - max upload rate (limiting max download rate) is  $u_s + \sum u_i$

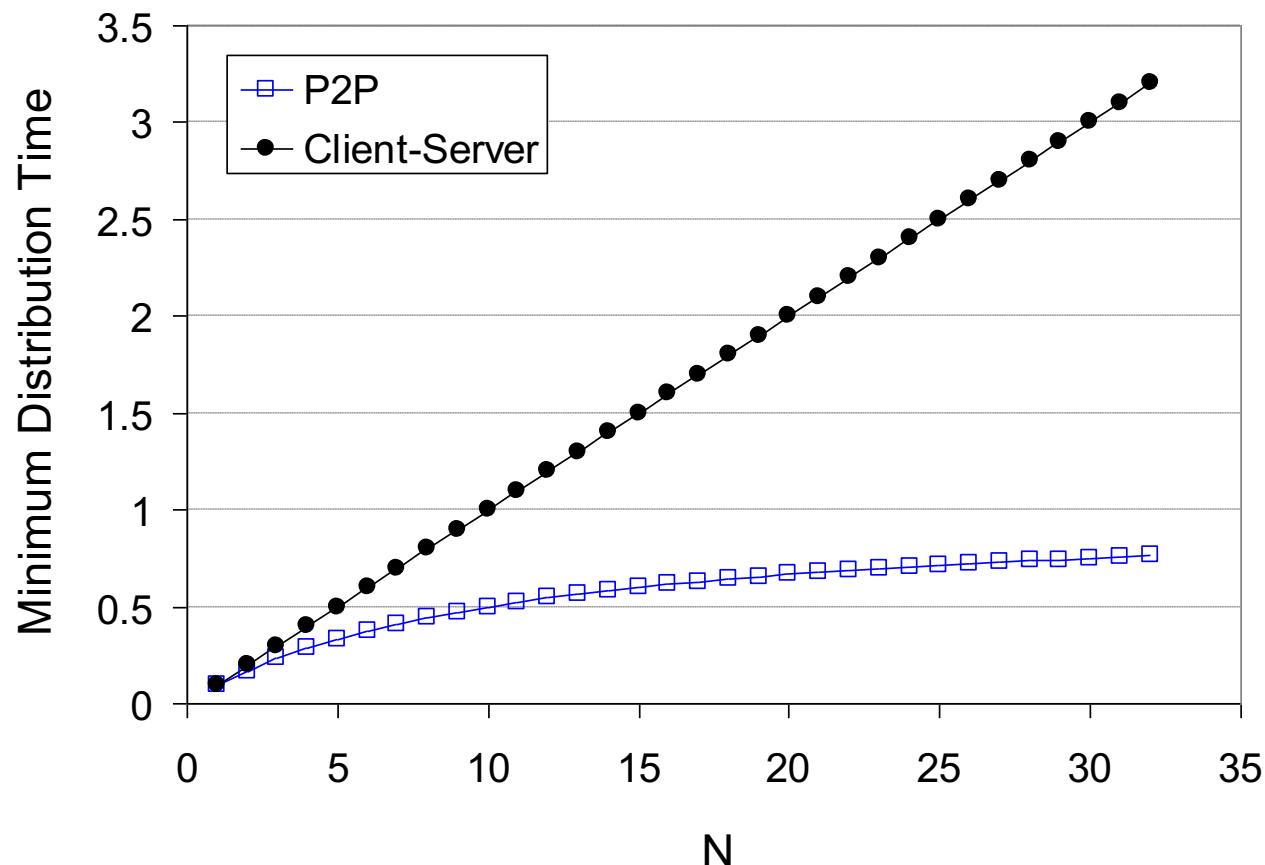
time to distribute  $F$   
to  $N$  clients using  
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in  $N$  ...  
... but so does this, as each peer brings service capacity

# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$

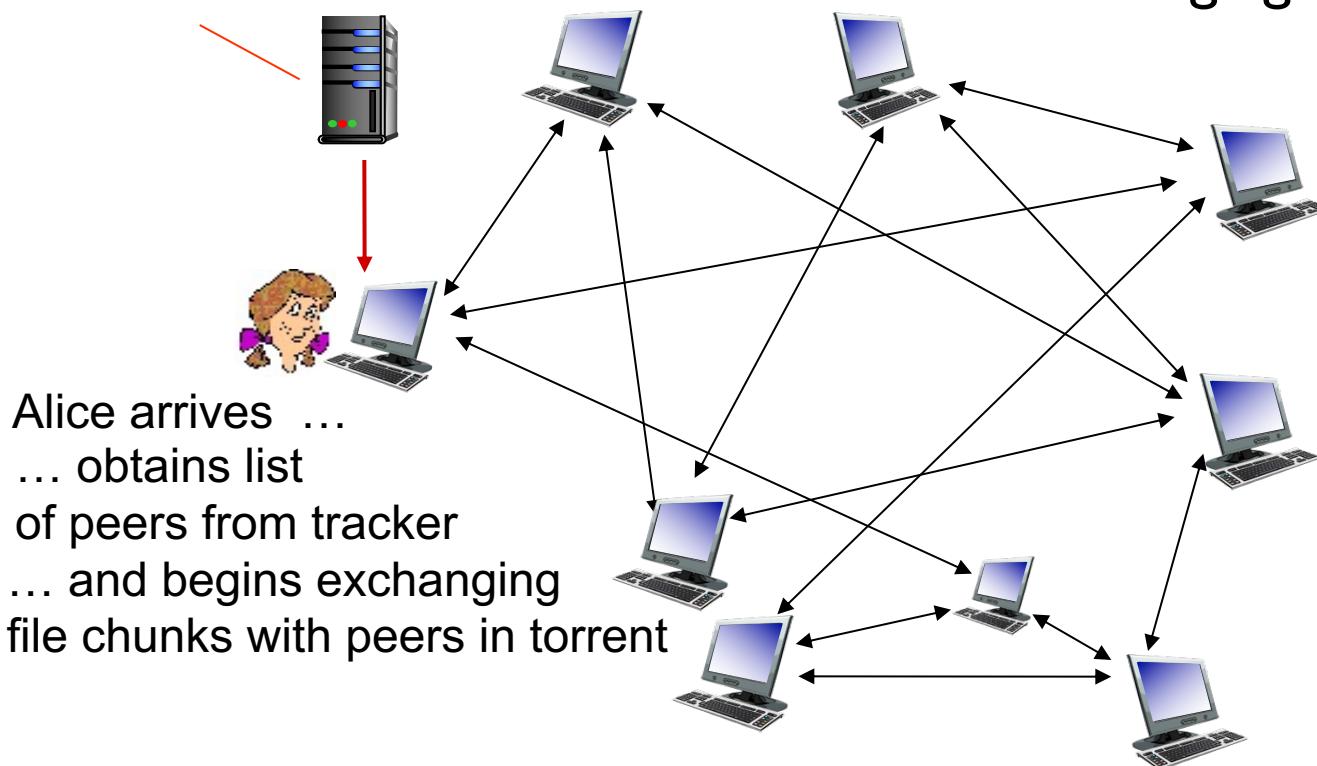


# P2P file distribution: BitTorrent

- ❖ file divided into 256Kb chunks
- ❖ peers in torrent send/receive file chunks

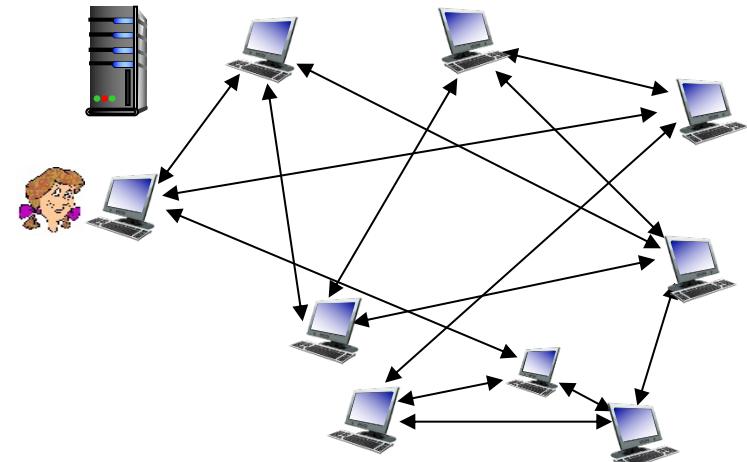
*tracker*: tracks peers  
participating in torrent

*torrent*: group of peers  
exchanging chunks of a file



# P2P file distribution: BitTorrent

- ❖ peer joining torrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- ❖ while downloading, peer uploads chunks to other peers
- ❖ peer may change peers with whom it exchanges chunks
- ❖ *churn*: peers may come and go
- ❖ once peer has entire file, it may (selfishly) leave or remain in torrent



# BitTorrent: requesting, sending file chunks

## *requesting chunks:*

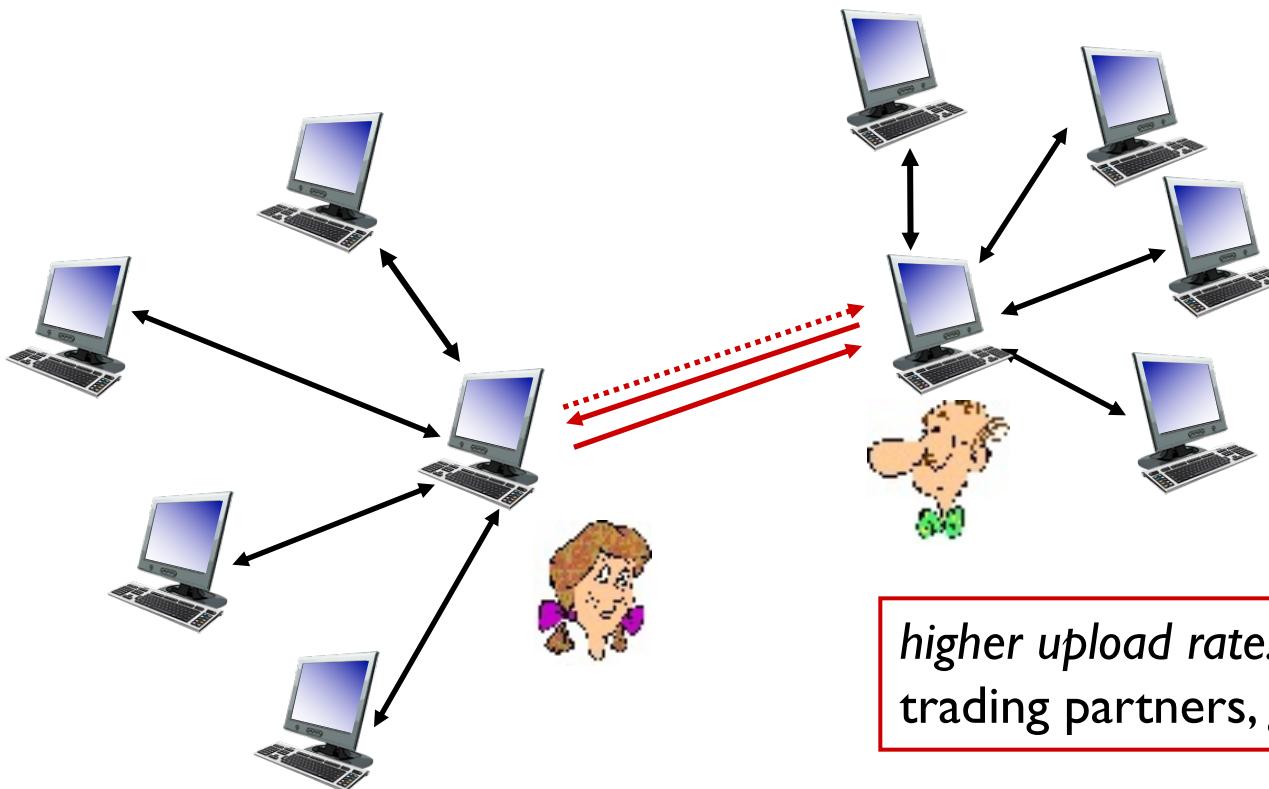
- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, Alice asks each peer for list of chunks that they have
- ❖ Alice requests missing chunks from peers, **rarest first**

## *sending chunks: tit-for-tat*

- ❖ Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate **top 4 every 10** secs. these four **peers** are said to be **unchoked**.
- ❖ every 30 secs: randomly select another peer, starts sending chunks
  - “optimistically unchoke”

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



*higher upload rate: find better trading partners, get file faster !*

# Distributed Hash Table (DHT)

- ❖ Hash table
- ❖ DHT paradigm
- ❖ Circular DHT and overlay networks
- ❖ Peer churn

# Simple Database

Simple database with **(key, value)** pairs:

- Value: human name; Key: social security #

Value	Key
John Washington	132-54-3570
Diana Louise Jones	761-55-3791
Xiaoming Liu	385-41-0902
Rakesh Gopal	441-89-1956
Linda Cohen	217-66-5609
.....	.....
Lisa Kobayashi	177-23-0199

# Hash Table

- More convenient to store and search on numerical representation of key
- key = hash(original key)

Value	Key	Original Key
John Washington	8962458	132-54-3570
Diana Louise Jones	7800356	761-55-3791
Xiaoming Liu	1567109	385-41-0902
Rakesh Gopal	2360012	441-89-1956
Linda Cohen	5430938	217-66-5609
.....	.....	.....
Lisa Kobayashi	9290124	177-23-0199

# Distributed Hash Table (DHT)

---

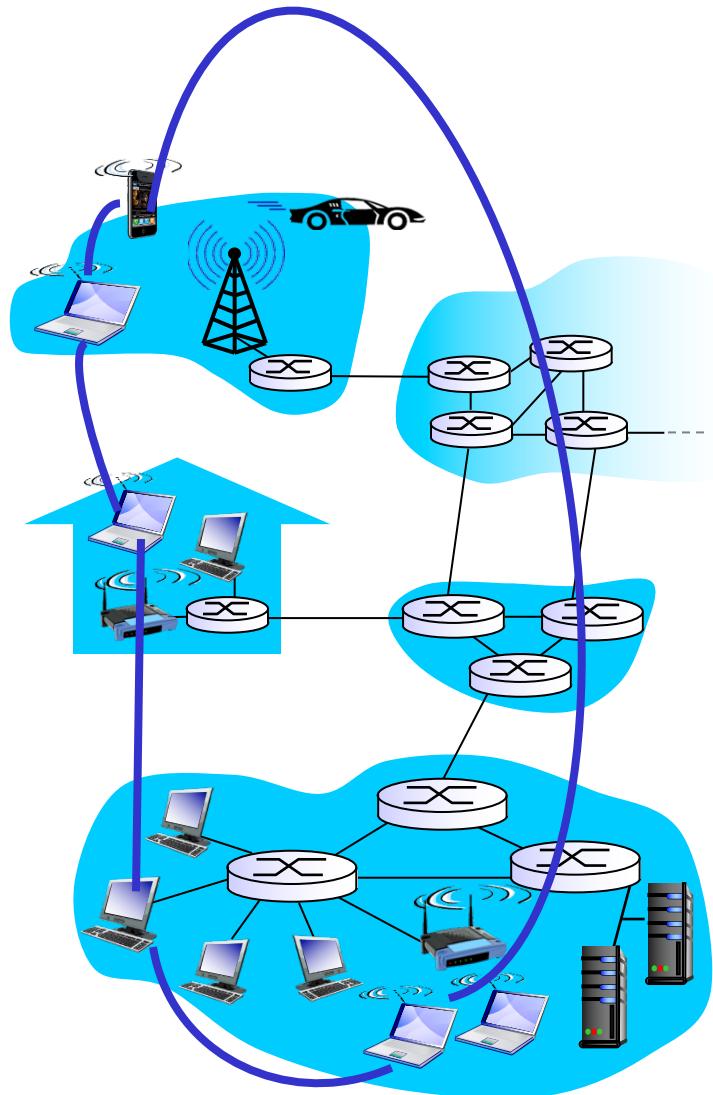
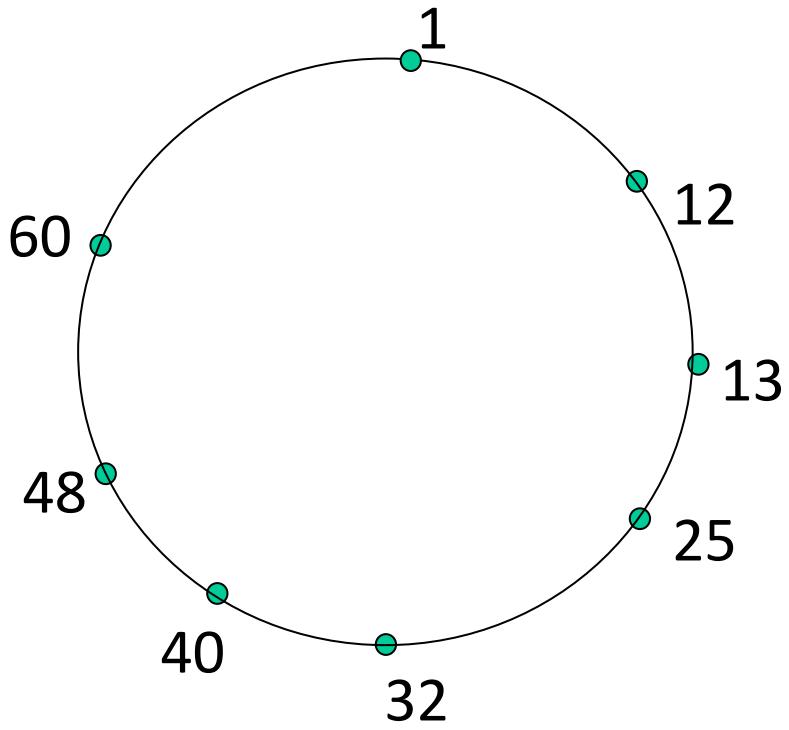
- ❖ Distribute (key, value) pairs over millions of peers
  - pairs are evenly distributed over peers
- ❖ Any peer can **query** database with a key
  - database returns value for the key
  - To resolve query, small number of messages exchanged among peers
- ❖ Each peer only knows about a small number of other peers
- ❖ Robust to peers coming and going (churn)

# Assign key-value pairs to peers

- ❖ rule: assign key-value pair to the peer that has the *closest* ID.
- ❖ convention: closest is the *immediate successor* of the key.
- ❖ e.g., ID space {0,1,2,3,...,63}
- ❖ suppose 8 peers: 1,12,13,25,32,40,48,60
  - If key = 51, then assigned to peer 60
  - If key = 60, then assigned to peer 60
  - If key = 61, then assigned to peer 1

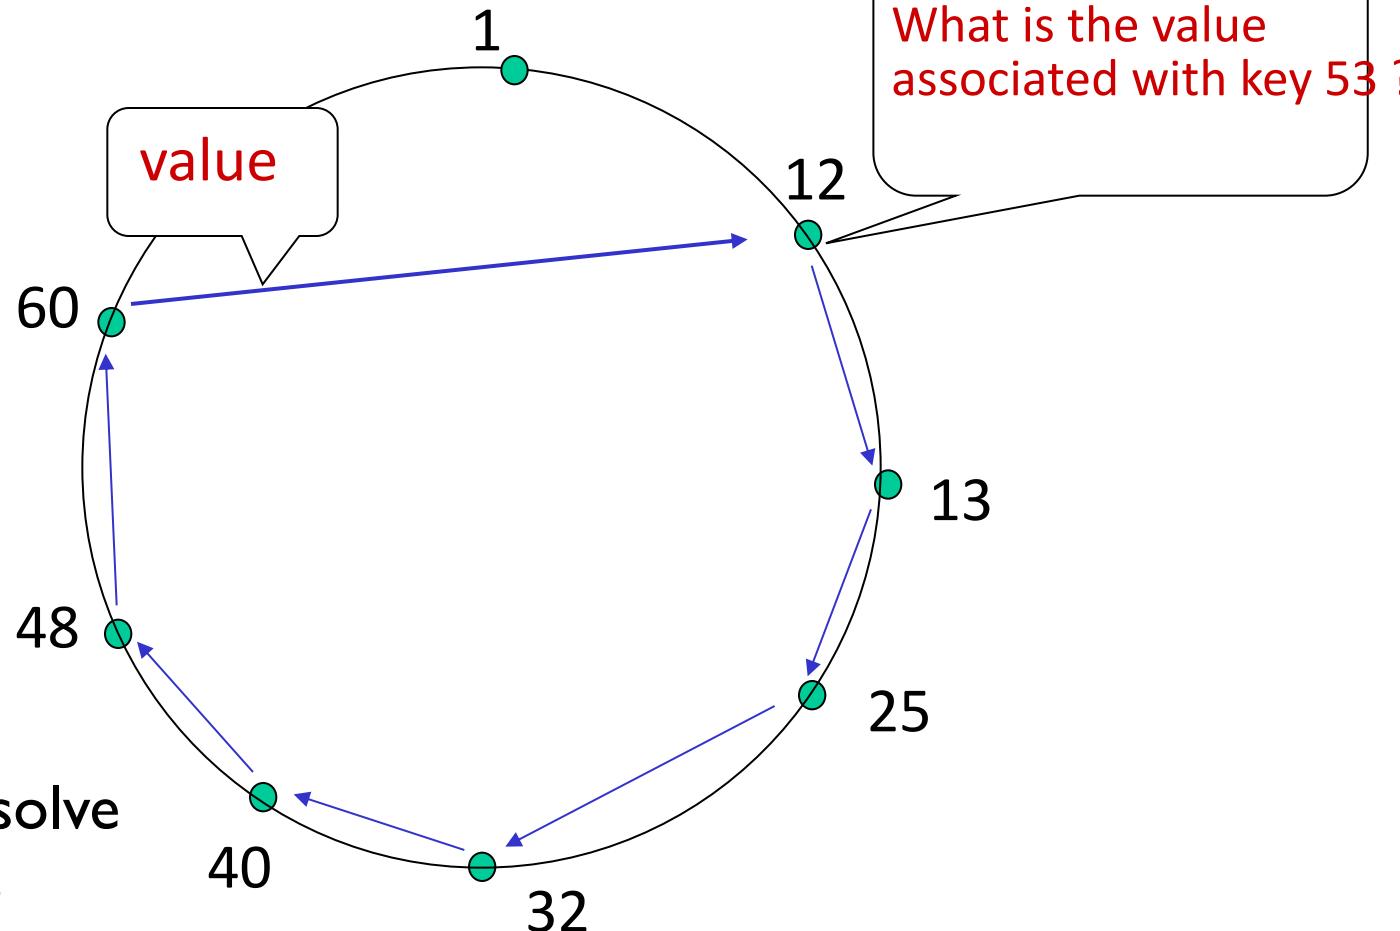
# Circular DHT

- each peer *only* aware of immediate successor and predecessor.



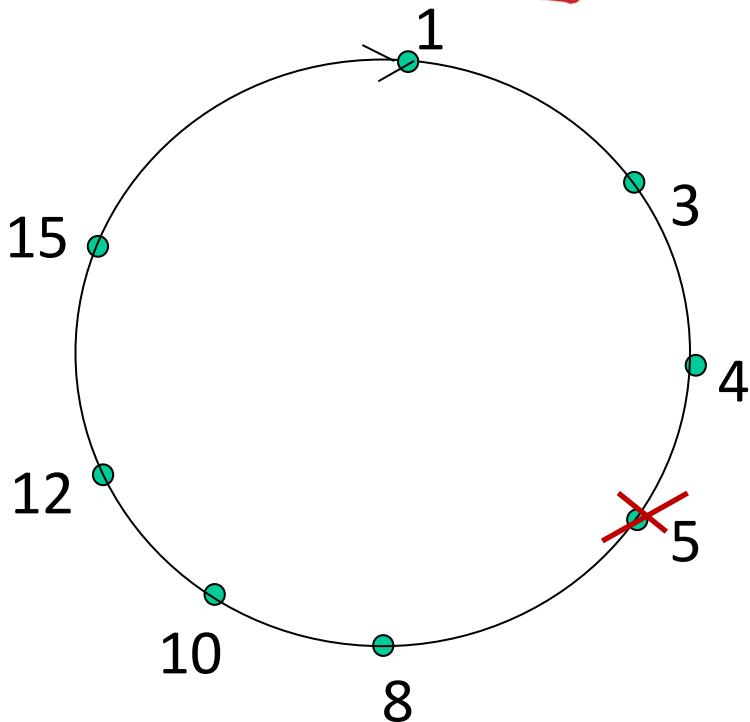
“overlay network”

# Resolving a query



$O(N)$  messages  
on average to resolve  
query, when there  
are  $N$  peers

# Peer churn

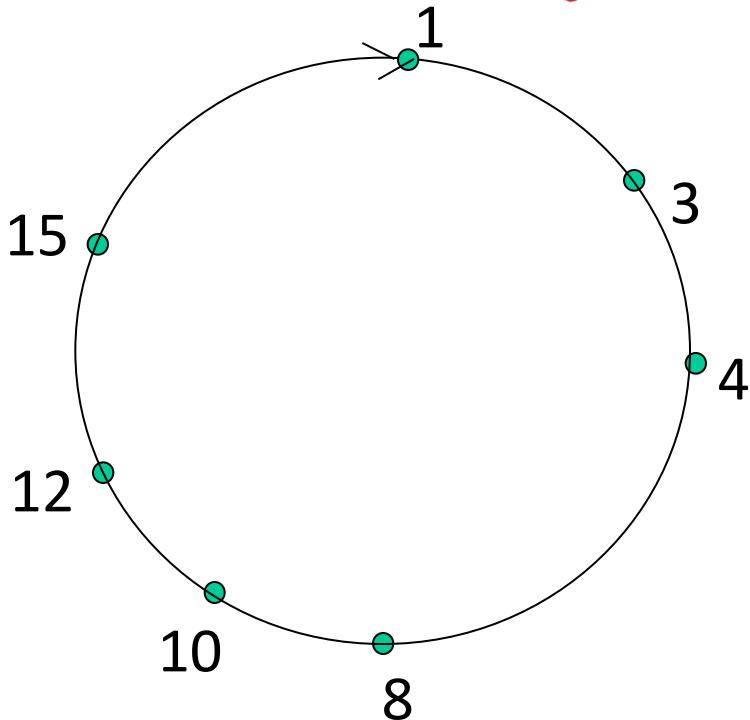


*example: peer 5 abruptly leaves*

## handling peer churn:

- ❖ peers may come and go (churn)
- ❖ each peer knows address of its two successors
- ❖ each peer periodically pings its two successors to check aliveness
- ❖ if immediate successor leaves, choose next successor as new immediate successor

# Peer churn



*example: peer 5 abruptly leaves*

- ❖ peer 4 detects peer 5's departure; makes 8 its immediate successor
- ❖ 4 asks 8 who its immediate successor is; makes 8's immediate successor its second successor.

**handling peer churn:**

- ❖ peers may come and go (churn)
- ❖ each peer knows address of its two successors
- ❖ each peer periodically pings its two successors to check aliveness
- ❖ if immediate successor leaves, choose next successor as new immediate successor

# Video Streaming and CDNs: context

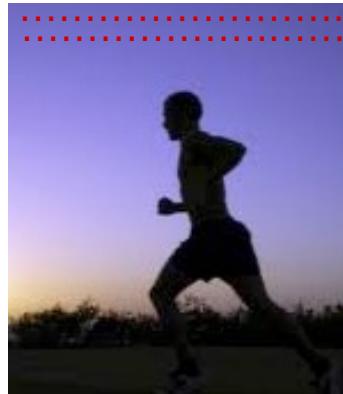
- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution:* distributed, application-level infrastructure



# Multimedia: video

- ❖ video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- ❖ digital image: array of pixels
  - each pixel represented by bits
- ❖ coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$

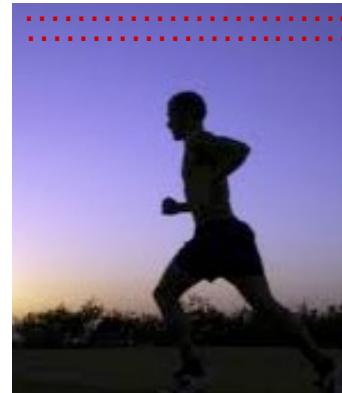


frame  $i+1$

# Multimedia: video

- **CBR: (constant bit rate):**  
video encoding rate fixed
- **VBR: (variable bit rate):**  
video encoding rate changes  
as amount of spatial,  
temporal coding changes
- **examples:**
  - MPEG I (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

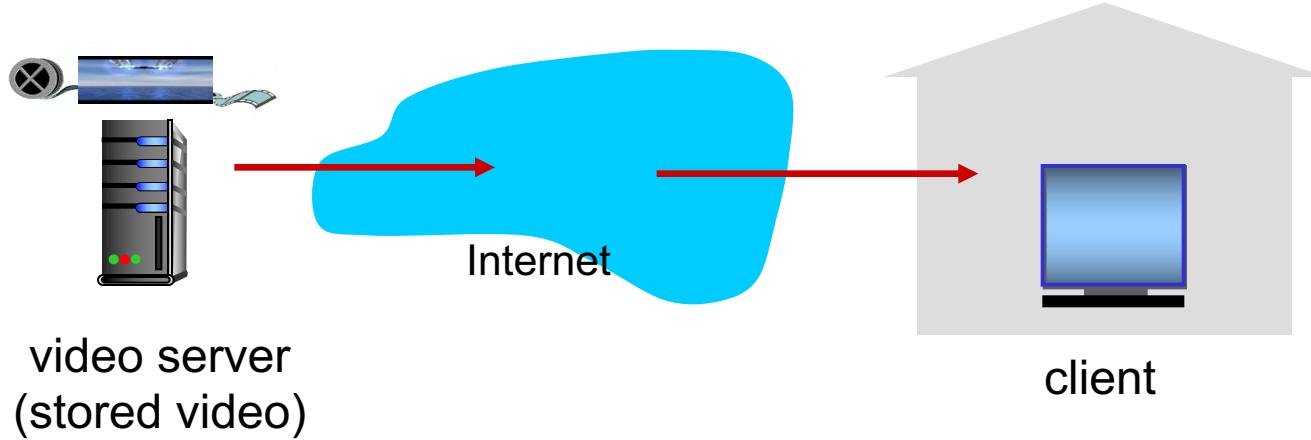
*temporal coding example:*  
instead of sending complete frame at  $i+1$ ,  
send only differences from frame  $i$



frame  $i+1$

# Streaming stored video:

simple scenario:



# Streaming multimedia: DASH

- ❖ *DASH: Dynamic, Adaptive Streaming over HTTP*
- ❖ *server:*
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
  - *manifest file:* provides URLs for different chunks
- ❖ *client:*
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates at different points in time (depending on available bandwidth at time)

# Streaming multimedia: DASH

- ❖ *DASH: Dynamic, Adaptive Streaming over HTTP*
- ❖ “*intelligence*” at client: client determines
  - *when* to request chunk (so that buffer starvation, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

# Content distribution networks

- ❖ *challenge*: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- ❖ *option 1*: single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long path to distant clients
  - multiple copies of video sent over outgoing link

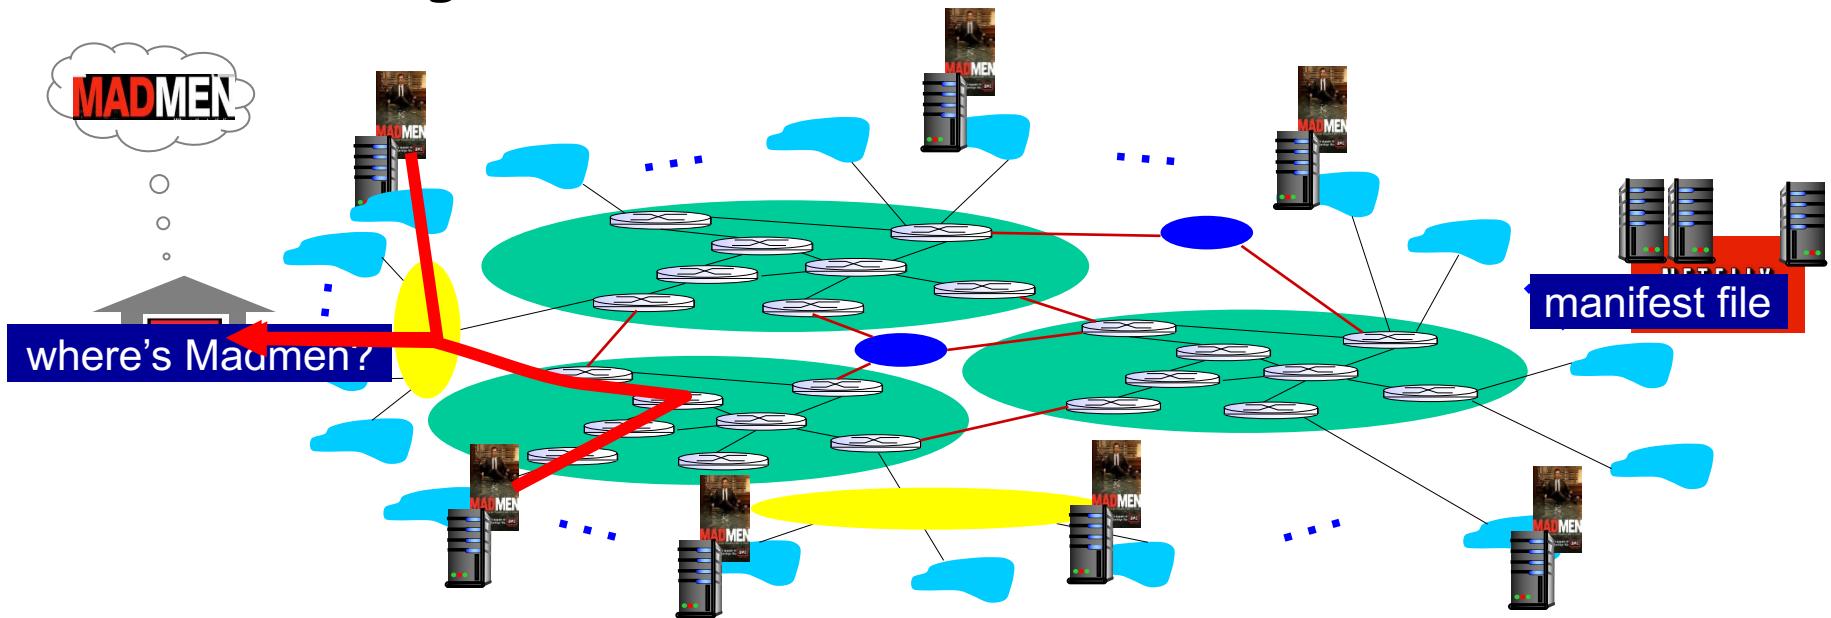
....quite simply: this solution *doesn't scale*

# Content distribution networks

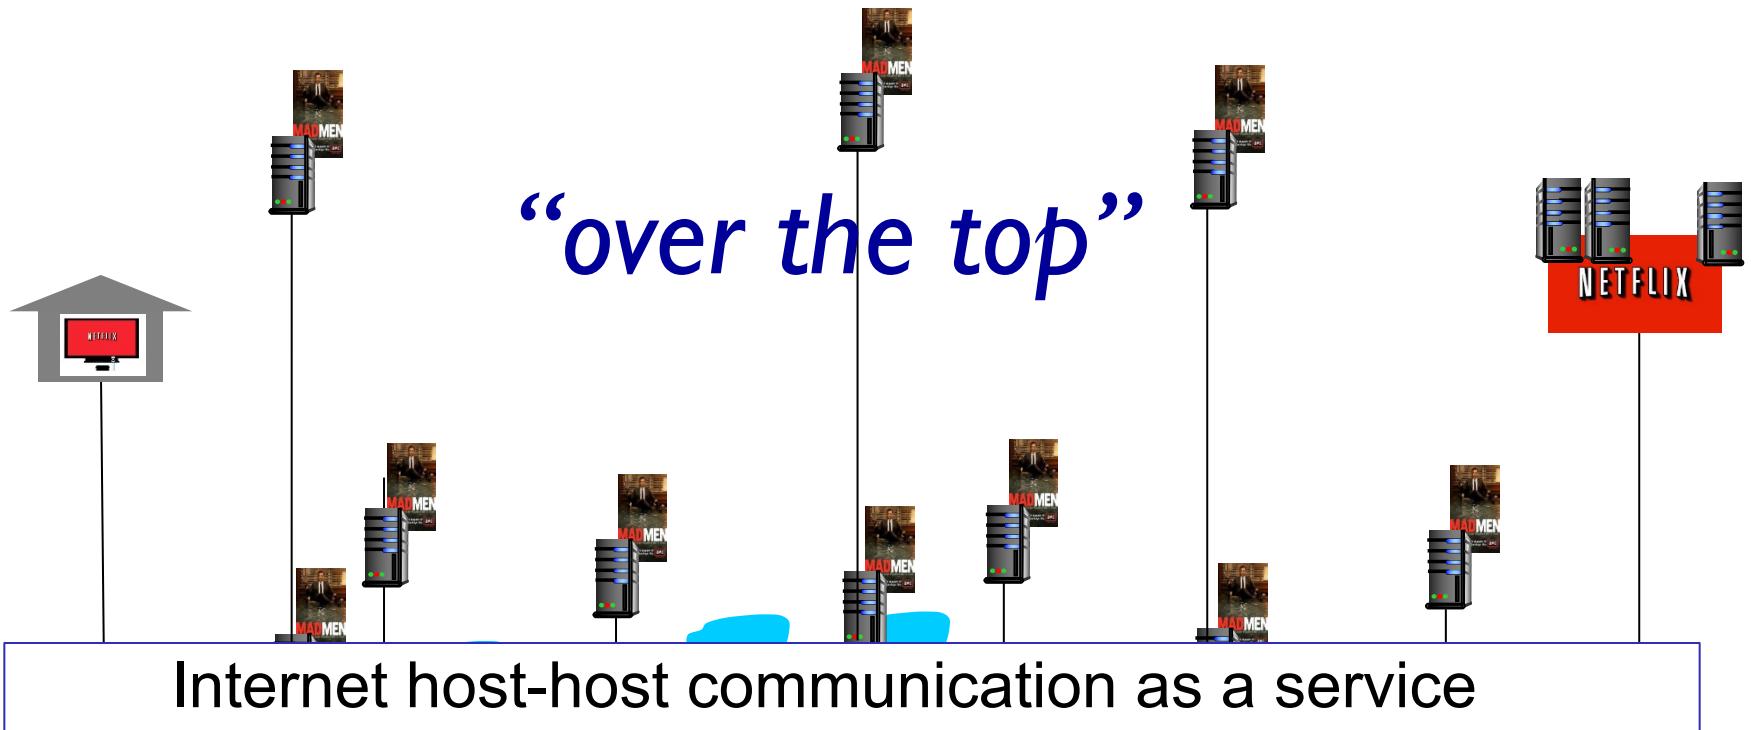
- ❖ ***challenge:*** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- ❖ ***option 2:*** store/serve multiple copies of videos at multiple geographically distributed sites (***CDN***)
  - ***enter deep:*** push CDN servers deep into many access networks
    - close to users
    - used by Akamai, 1700 locations
  - ***bring home:*** smaller number (10's) of larger clusters in POPs near (but not within) access networks
    - used by Limelight
  - PoP : Point of Presence

# Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



# Content Distribution Networks (CDNs)



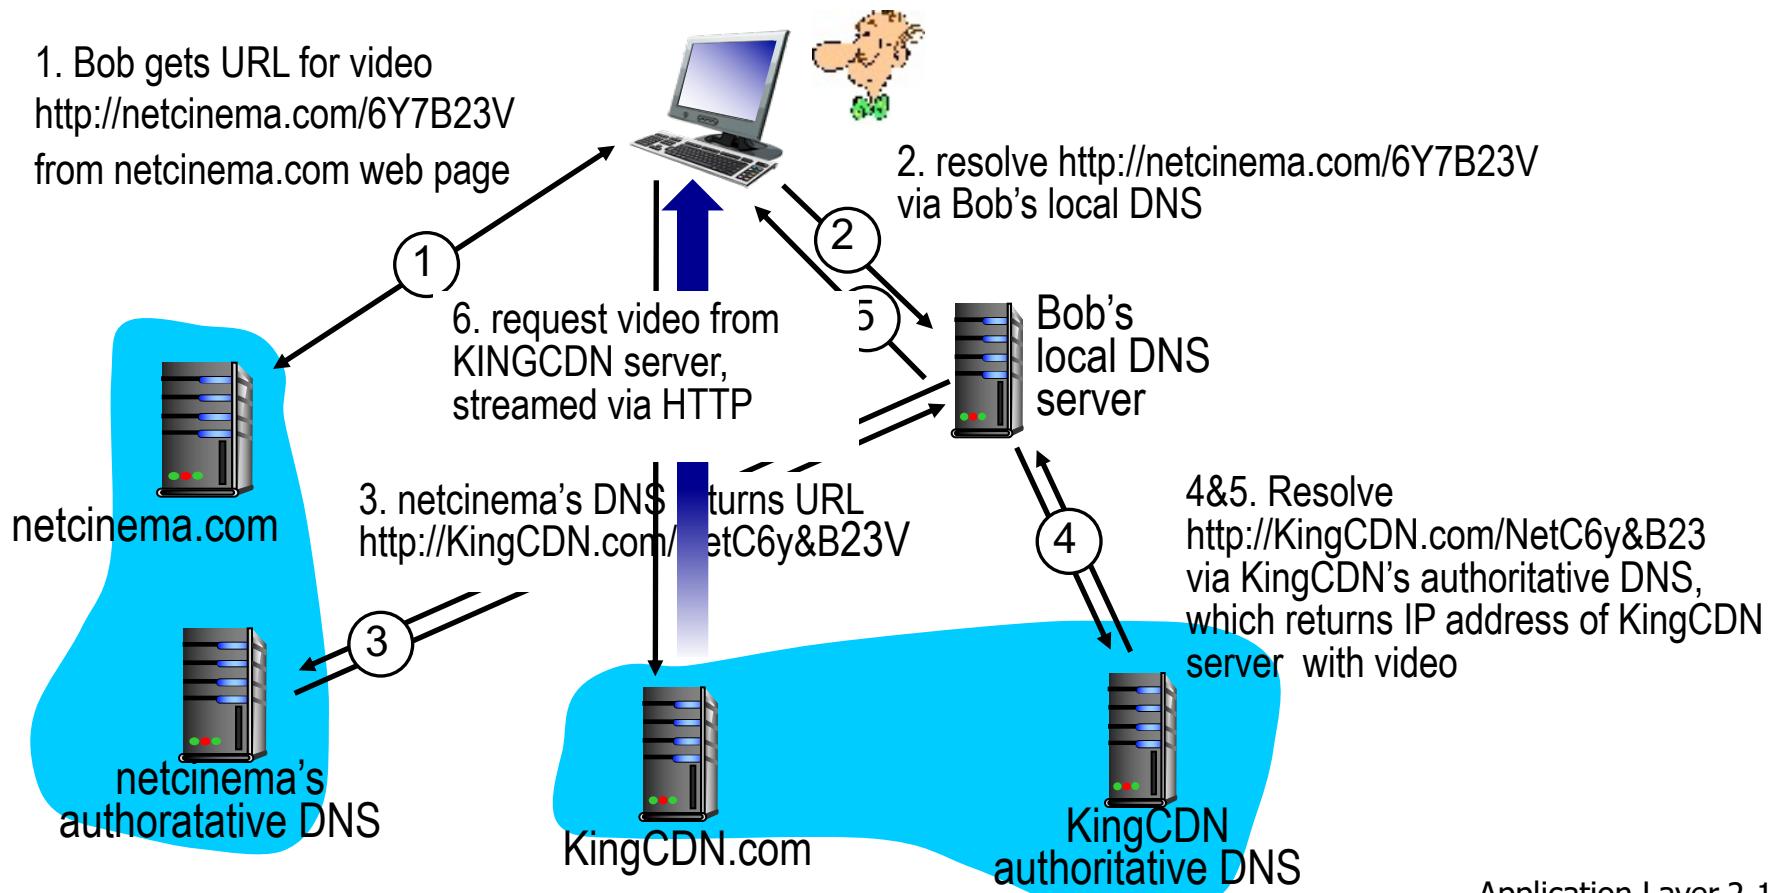
*Over the Top(OTT) challenges:* coping with a congested Internet

- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?  
more .. in chapter 7

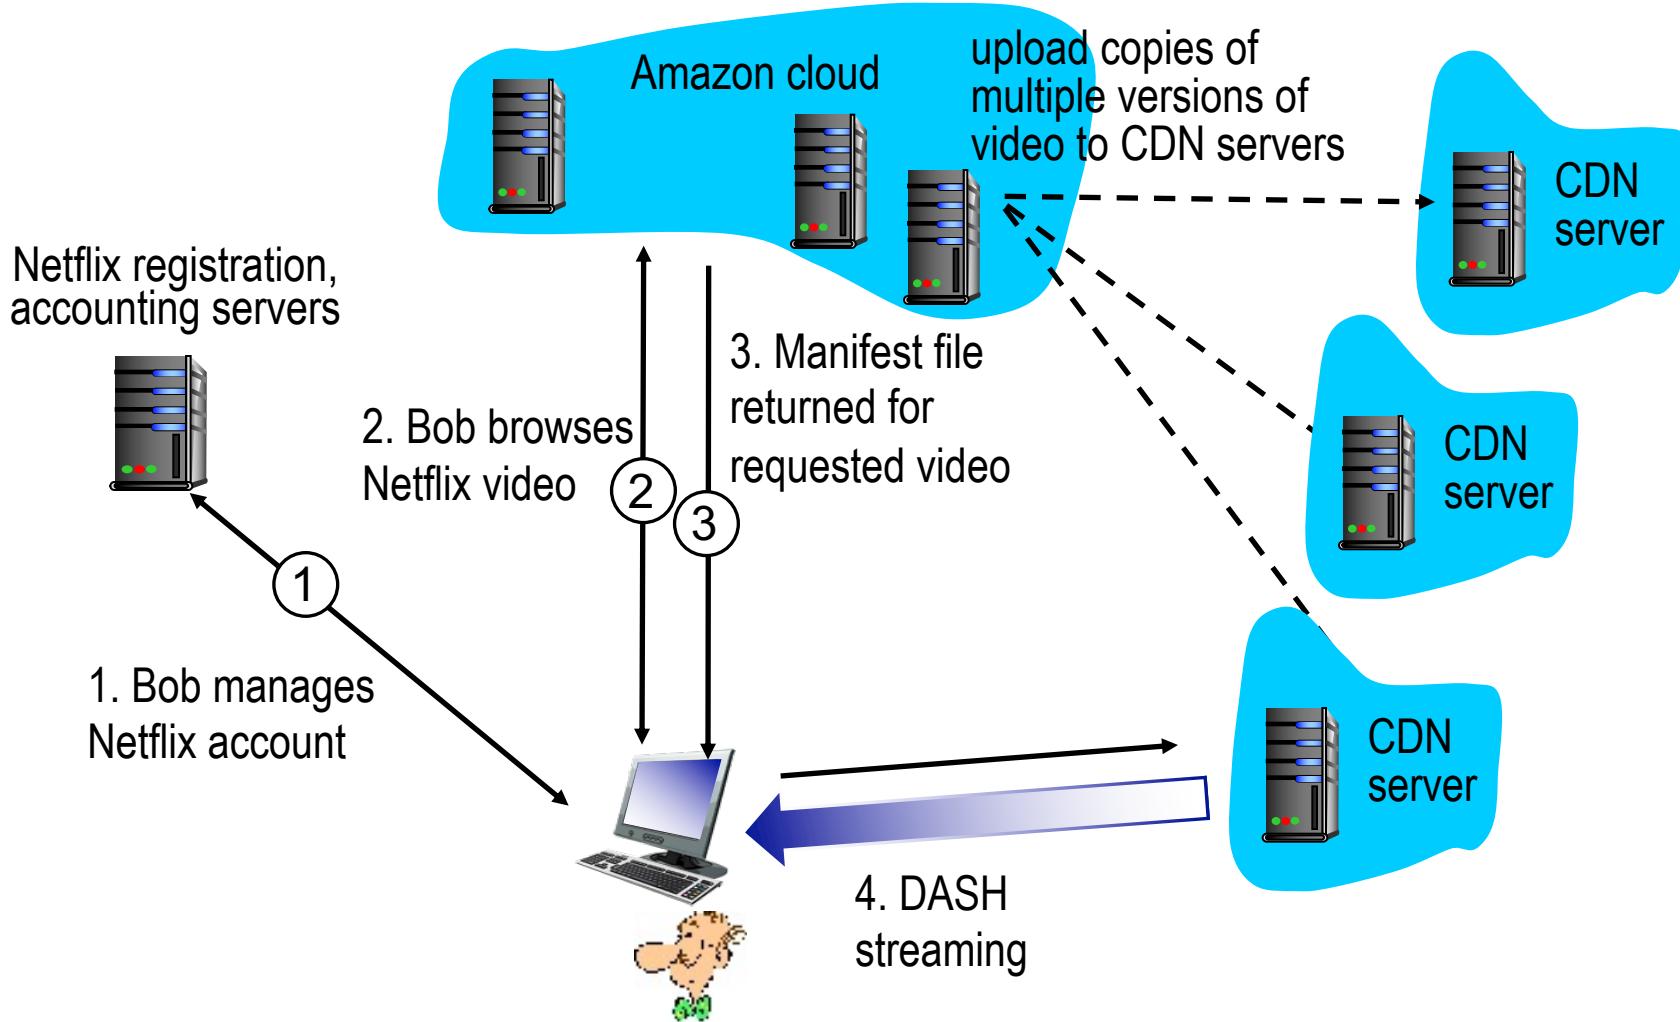
# CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



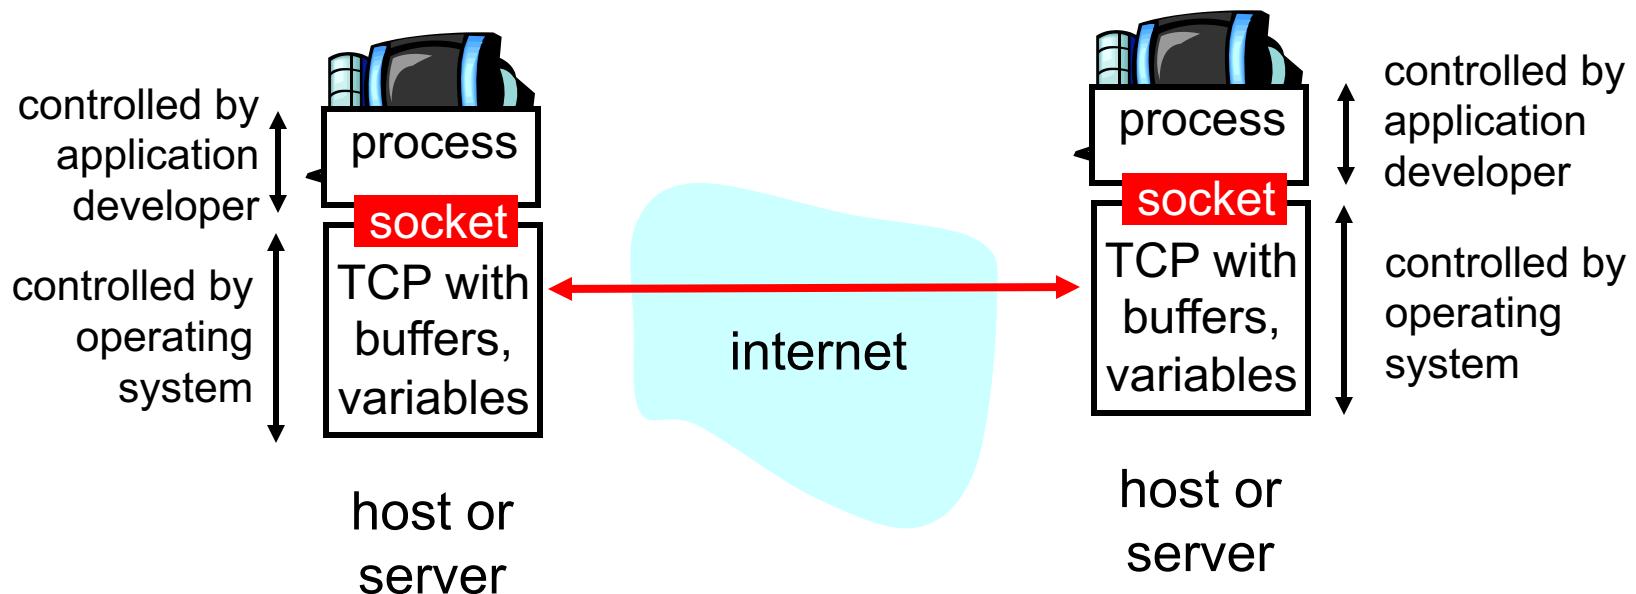
# Case study: Netflix



# Socket-programming using TCP

Socket: a door between application process and end-end-transport protocol (UCP or TCP)

TCP service: reliable transfer of **bytes** from one process to another



# Socket programming *with TCP*

## Client must contact server

- ❖ server process must first be running
- ❖ server must have created socket (door) that welcomes client's contact

## Client contacts server by:

- ❖ creating client-local TCP socket
- ❖ specifying IP address, port number of server process
- ❖ When **client creates socket**: client TCP establishes connection to server TCP

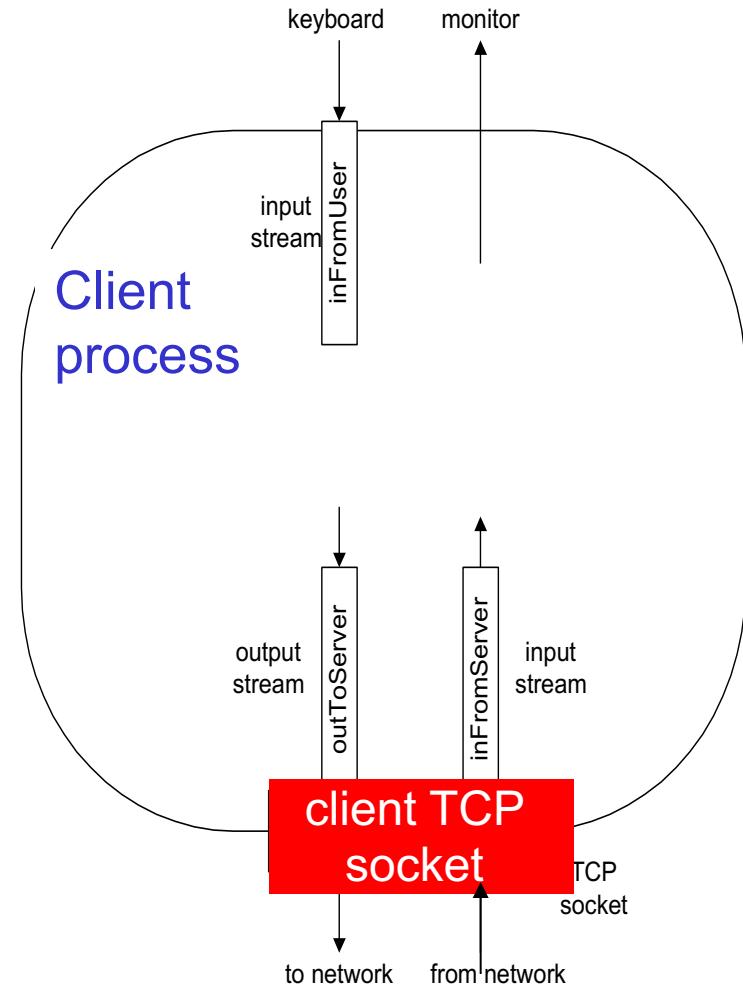
- ❖ When contacted by client, **server TCP creates new socket** for server process to communicate with client
  - allows server to talk with multiple clients
  - source port numbers used to distinguish clients (**more in Chap 3**)

application viewpoint

*TCP provides reliable, in-order transfer of bytes (“pipe”) between client and server*

# Stream jargon

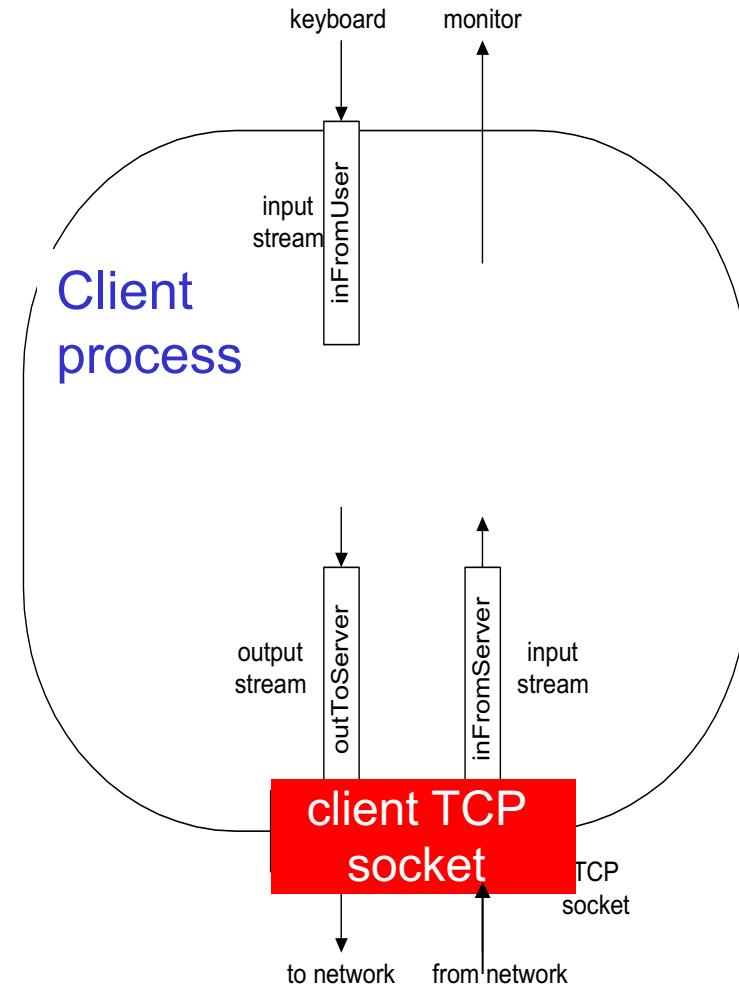
- ❖ A **stream** is a sequence of characters that flow into or out of a process.
- ❖ An **input stream** is attached to some input source for the process, eg, keyboard or socket.
- ❖ An **output stream** is attached to an output source, eg, monitor or socket.



# Socket programming with TCP

## Example client-server app:

- 1) client reads line from standard input (**inFromUser** stream), sends to server via socket (**outToServer** stream)
- 2) server reads line from socket
- 3) server converts line to uppercase, sends back to client
- 4) client reads, prints modified line from socket (**inFromServer** stream)



# Client/server socket interaction: TCP

Server (running on `hostid`)

Client

```
create socket,  
port=x, for  
incoming request:  
welcomeSocket =  
    ServerSocket()
```

```
wait for incoming  
connection request  
connectionSocket =  
    welcomeSocket.accept()
```

```
read request from  
connectionSocket
```

```
write reply to  
connectionSocket
```

```
close  
connectionSocket
```

TCP  
connection setup

```
create socket,  
connect to hostid, port=x  
clientSocket =  
    Socket()
```

```
send request using  
clientSocket
```

```
read reply from  
clientSocket
```

```
close  
clientSocket
```

# Example app: TCP client

## *Python TCPClient*

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```

create TCP socket for  
server, remote port 12000



clientSocket = socket(AF\_INET, SOCK\_STREAM)



No need to attach server  
name, port



clientSocket.connect((serverName,serverPort))

# Example app: TCP server

## *Python TCPServer*

create TCP welcoming  
socket

```
from socket import *
serverPort = 12000
```

server begins listening for  
incoming TCP requests

```
serverSocket = socket(AF_INET,SOCK_STREAM)
serverSocket.bind(("",serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
```

loop forever

```
while True:
    connectionSocket, addr = serverSocket.accept()
```

server waits on accept()  
for incoming requests, new  
socket created on return

```
sentence = connectionSocket.recv(1024).decode()
```

```
capitalizedSentence = sentence.upper()
```

```
connectionSocket.send(capitalizedSentence.
                      encode())
```

read bytes from socket (but  
not address as in UDP)

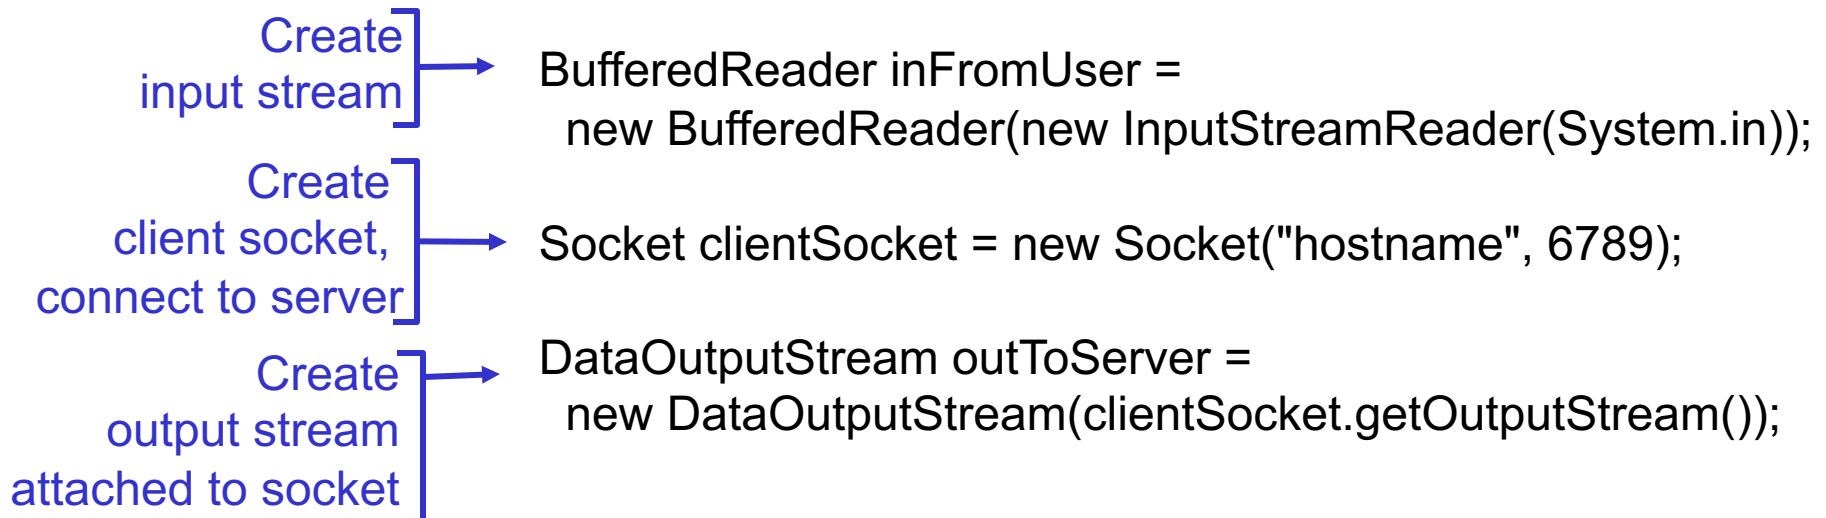
```
connectionSocket.close()
```

close connection to this  
client (but *not* welcoming  
socket)

# Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPClient {

    public static void main(String argv[]) throws Exception
    {
        String sentence;
        String modifiedSentence;


        BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));

        Socket clientSocket = new Socket("hostname", 6789);

        DataOutputStream outToServer =
            new DataOutputStream(clientSocket.getOutputStream());
    }
}
```

# Example: Java client (TCP), cont.

```
        Create  
        input stream  
attached to socket }→ BufferedReader inFromServer =  
new BufferedReader(new  
InputStreamReader(clientSocket.getInputStream()));  
  
Send line }→ to server sentence = inFromUser.readLine();  
outToServer.writeBytes(sentence + '\n');  
  
Read line }→ from server modifiedSentence = inFromServer.readLine();  
System.out.println("FROM SERVER: " + modifiedSentence);  
  
clientSocket.close();  
  
}  
}
```

# Example: Java server (TCP)

```
import java.io.*;
import java.net.*;

class TCPServer {

    public static void main(String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;

        ServerSocket welcomeSocket = new ServerSocket(6789);

        while(true) {
            Socket connectionSocket = welcomeSocket.accept();

            BufferedReader inFromClient =
                new BufferedReader(new
                    InputStreamReader(connectionSocket.getInputStream()));

            // Create output stream, attached to socket
            DataOutputStream outToClient =
                new DataOutputStream(connectionSocket.getOutputStream());
        }
    }
}
```

**Create welcoming socket at port 6789**

**Wait, on welcoming socket for contact by client**

**Create input stream, attached to socket**

# Example: Java server (TCP), cont

The diagram illustrates a Java code snippet for a TCP server with annotations explaining its functionality:

```
DataOutputStream outToClient =  
    new DataOutputStream(connectionSocket.getOutputStream());  
  
clientSentence = inFromClient.readLine();  
  
capitalizedSentence = clientSentence.toUpperCase() + '\n';  
  
outToClient.writeBytes(capitalizedSentence);
```

The annotations are:

- Create output stream, attached to socket → Points to the line: `DataOutputStream outToClient = new DataOutputStream(connectionSocket.getOutputStream());`
- Read in line from socket → Points to the line: `clientSentence = inFromClient.readLine();`
- Write out line to socket → Points to the line: `outToClient.writeBytes(capitalizedSentence);`
- End of while loop, loop back and wait for another client connection → Points to the closing brace } at the bottom of the code block.

# Socket programming *with UDP*

UDP: no “connection” between client and server

- ❖ no handshaking
- ❖ sender explicitly attaches IP address and port of destination to each packet
- ❖ server must extract IP address, port of sender from received packet

UDP: transmitted data may be received out of order, or lost

application viewpoint

*UDP provides unreliable transfer of groups of bytes (“datagrams”) between client and server*

# Client/server socket interaction: UDP

Server (running on `hostid`)

create socket,  
`port=x`, for  
incoming request:  
`serverSocket =  
DatagramSocket()`

read request from  
`serverSocket`  
  
write reply to  
`serverSocket`  
specifying client  
host address,  
port number

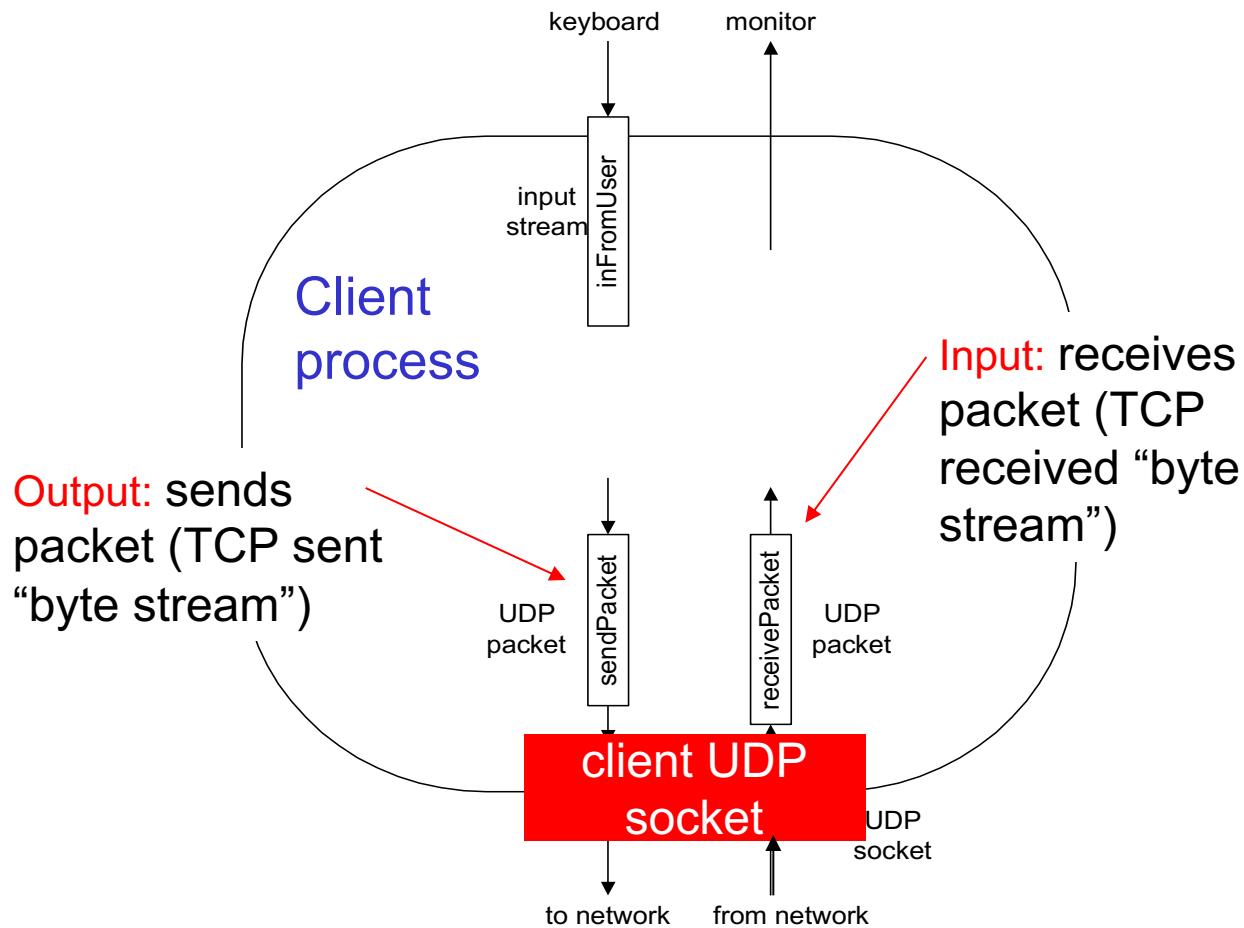
Client

create socket,  
`clientSocket =  
DatagramSocket()`

Create, address (`hostid, port=x`),  
send datagram request  
using `clientSocket`

read reply from  
`clientSocket`  
  
close  
`clientSocket`

# Example: Java client (UDP)



# Example: Java client (UDP)

```
import java.io.*;
import java.net.*;

class UDPClient {
    public static void main(String args[]) throws Exception
    {
        BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));

        DatagramSocket clientSocket = new DatagramSocket();

        InetAddress IPAddress = InetAddress.getByName("hostname");

        byte[] sendData = new byte[1024];
        byte[] receiveData = new byte[1024];

        String sentence = inFromUser.readLine();
        sendData = sentence.getBytes();
    }
}
```

**Create input stream** → BufferedReader inFromUser =  
new BufferedReader(new InputStreamReader(System.in));

**Create client socket** → DatagramSocket clientSocket = new DatagramSocket();

**Translate hostname to IP address using DNS** → InetAddress IPAddress = InetAddress.getByName("hostname");

# Example: Java client (UDP), cont.

Create datagram with

  data-to-send,

length, IP addr, port

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, IPAddress, 9876);
```

Send datagram  
to server

```
clientSocket.send(sendPacket);
```

```
DatagramPacket receivePacket =  
    new DatagramPacket(receiveData, receiveData.length);
```

Read datagram  
from server

```
clientSocket.receive(receivePacket);
```

```
String modifiedSentence =  
    new String(receivePacket.getData());
```

```
System.out.println("FROM SERVER:" + modifiedSentence);  
clientSocket.close();
```

```
}
```

```
}
```

# Example: Java server (UDP)

```
import java.io.*;
import java.net.*;

class UDPServer {
    public static void main(String args[]) throws Exception
    {
        DatagramSocket serverSocket = new DatagramSocket(9876);

        byte[] receiveData = new byte[1024];
        byte[] sendData = new byte[1024];

        while(true)
        {
            DatagramPacket receivePacket =
                new DatagramPacket(receiveData, receiveData.length);
            serverSocket.receive(receivePacket);
        }
    }
}
```

**Create datagram socket at port 9876**

**Create space for received datagram**

**Receive datagram**

# Example: Java server (UDP), cont

```
String sentence = new String(receivePacket.getData());  
Get IP addr  
port #, of  
sender  
→ InetAddress IPAddress = receivePacket.getAddress();  
→ int port = receivePacket.getPort();  
  
String capitalizedSentence = sentence.toUpperCase();  
  
sendData = capitalizedSentence.getBytes();  
Create datagram  
to send to client  
→ DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, IPAddress,  
                       port);  
  
Write out  
datagram  
to socket  
→ serverSocket.send(sendPacket);  
}  
}  
} }  
End of while loop,  
loop back and wait for  
another datagram
```

# Building a simple Web server

- ❖ handles one HTTP request
- ❖ accepts the request
- ❖ parses header
- ❖ obtains requested file from server's file system
- ❖ creates HTTP response message:
  - header lines + file
- ❖ sends response to client
- ❖ after creating server, you can request file using a browser (eg IE explorer)
- ❖ see text for details

# Chapter 2: summary

*our study of network apps now complete!*

- ❖ application architectures
  - client-server
  - P2P
- ❖ application service requirements:
  - reliability, bandwidth, delay
- ❖ Internet transport service model
  - connection-oriented, reliable: TCP
  - unreliable, datagrams: UDP
- ❖ specific protocols:
  - HTTP
  - FTP
  - SMTP, POP, IMAP
  - DNS
  - P2P: BitTorrent, DHT

# Chapter 2: summary

*most importantly: learned about protocols!*

- ❖ typical request/reply message exchange:
  - client requests info or service
  - server responds with data, status code
- ❖ message formats:
  - headers: fields giving info about data
  - data: info being communicated

*important themes:*

- ❖ control vs. data msgs
  - in-band, out-of-band
- ❖ centralized vs. decentralized
- ❖ stateless vs. stateful
- ❖ reliable vs. unreliable msg transfer
- ❖ “complexity at network edge”